

CME 295: Transformers & Large Language Models



Afshin Amidi & Shervine Amidi



Teaching staff



Afshine and Shervine

Teaching staff



Afshine
Centrale Paris ('16), MIT ('17)
Uber, Google, Netflix



Shervine
Centrale Paris ('16), Stanford ('19)
Uber, Google, Netflix

Welcome to CME 295!

Goals.

1. Understand how **Transformers** work and how they **relate** to **LLMs**
2. Learn how **LLMs** are **trained** and used in various **applications**

Welcome to CME 295!

Goals.

1. Understand how **Transformers** work and how they **relate** to **LLMs**
2. Learn how **LLMs** are **trained** and used in various **applications**

Audience.

- Interested in LLMs
 - Career goal
 - Personal project
 - "AI literacy" or curiosity
- Prerequisite: machine learning basics, linear algebra

Logistics

Date & time.

- Fridays from 3:30pm to 5:20pm
- Thornton 110

Logistics

Date & time.

- Fridays from 3:30pm to 5:20pm
- Thornton 110

Details about the class.

- 2 units
- Letter or Credit/No credit
- Lectures are recorded
- Midterm (50% grade), scheduled on October 24th
- Final exam (50% grade), week of December 8th (exact date TBD)

Material

Class website. cme295.stanford.edu

- Contains syllabus & logistics
- Slides and recordings will be posted there

Material

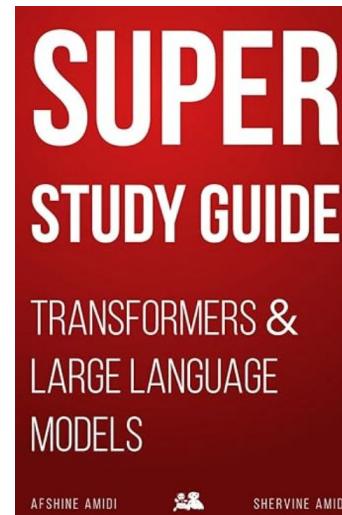
Class website. cme295.stanford.edu

- Contains syllabus & logistics
- Slides and recordings will be posted there

Class textbook.

Super Study Guide:
Transformers &
Large Language Models

Book: <https://superstudy.guide>



Material

CME 295 – TRANSFORMERS & LARGE LANGUAGE MODELS

<https://cme295.stanford.edu>

VIP Cheatsheet:

Transformers & Large Language Models

Afshin AMIDI and Shervine AMIDI

March 23, 2025

This VIP cheatsheet gives an overview of what is in the 'Super Study Guide: Transformers & Large Language Models' book, which contains ~600 illustrations over 250 pages and goes into the following concepts in depth. You can find more details at <https://superstudy.guide>.

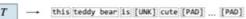
1 Foundations

1.1 Tokens

□ Definition – A token is an indivisible unit of text, such as a word, subword or character, and is part of a predefined vocabulary.

Remark: The unknown token `[UNK]` represents unknown pieces of text while the padding token `[PAD]` is used to fill empty positions to ensure consistent input sequence lengths.

□ Tokenizer – A tokenizer T divides text into tokens of an arbitrary level of granularity.

this teddy bear is reaaaally cute → 

Here are the main types of tokenizers:

Type	Pros	Cons	Illustration
Word	<ul style="list-style-type: none">Easy to interpretShort sequence	<ul style="list-style-type: none">Large vocabulary sizeWord variations not handled	
Subword	<ul style="list-style-type: none">Word roots leveragedIntuitive embeddings	<ul style="list-style-type: none">Increased sequence lengthTokenization more complex	
Character	<ul style="list-style-type: none">No out-of-vocabulary concernsSmall vocabulary size	<ul style="list-style-type: none">Much longer sequence lengthPatterns hard to interpret because too low-level	
Byte			

Remark: Byte-Pair Encoding (BPE) and Unigram are commonly-used subword-level tokenizers.

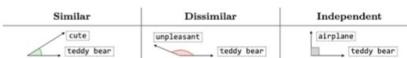
1.2 Embeddings

□ Definition – An embedding is a numerical representation of an element (e.g. token, sentence) and is characterized by a vector $\mathbf{x} \in \mathbb{R}^n$.

□ Similarity – The cosine similarity between two tokens t_1, t_2 is quantified by:

$$\text{similarity}(t_1, t_2) = \frac{t_1 \cdot t_2}{\|t_1\| \|t_2\|} = \cos(\theta) \in [-1, 1]$$

The angle θ characterizes the similarity between the two tokens:



Remark: Approximate Nearest Neighbors (ANN) and Locality Sensitive Hashing (LSH) are methods that approximate the similarity operation efficiently over large databases.

2 Transformers

2.1 Attention

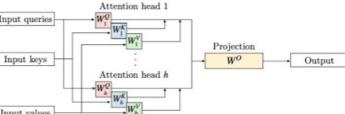
□ Formula – Given a query q , we want to know which key k the query should pay 'attention' to with respect to the associated value v .



Attention can be efficiently computed using matrices Q, K, V that contain queries q , keys k and values v respectively, along with the dimension d_k of keys:

$$\text{attention} = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

□ MHA – A Multi-Head Attention (MHA) layer performs attention computations across multiple heads, then projects the result in the output space.



It is composed of h attention heads as well as matrices W^Q, W^K, W^V that project the input to obtain query Q , keys K and values V . The projection is done using matrix W^O .

Remark: Grouped-Query Attention (GQA) and Multi-Query Attention (MQA) are variations of MHA that reduce computational overhead by sharing keys and values across attention heads.

2.2 Architecture

□ Overview – Transformer is a landmark model relying on the self-attention mechanism and is composed of encoders and decoders. Encoders compute meaningful embeddings of the input that are then used by decoders to predict the next token in the sequence.

Link to PDF: <https://github.com/afshinea/stanford-cme-295-transformers-large-language-models>

Class communications

Canvas.

- Announcements
- Class discussions *via Ed*

Class communications

Canvas.

- Announcements
- Class discussions via Ed

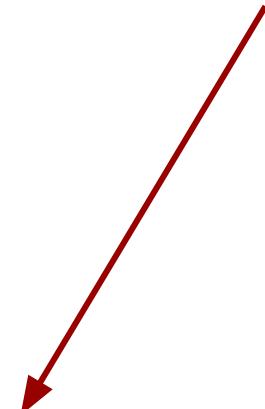
Any other general **inquiries / questions**, email us:

- cme295-aut2526-staff@lists.stanford.edu
- afshine@stanford.edu, shervine@stanford.edu

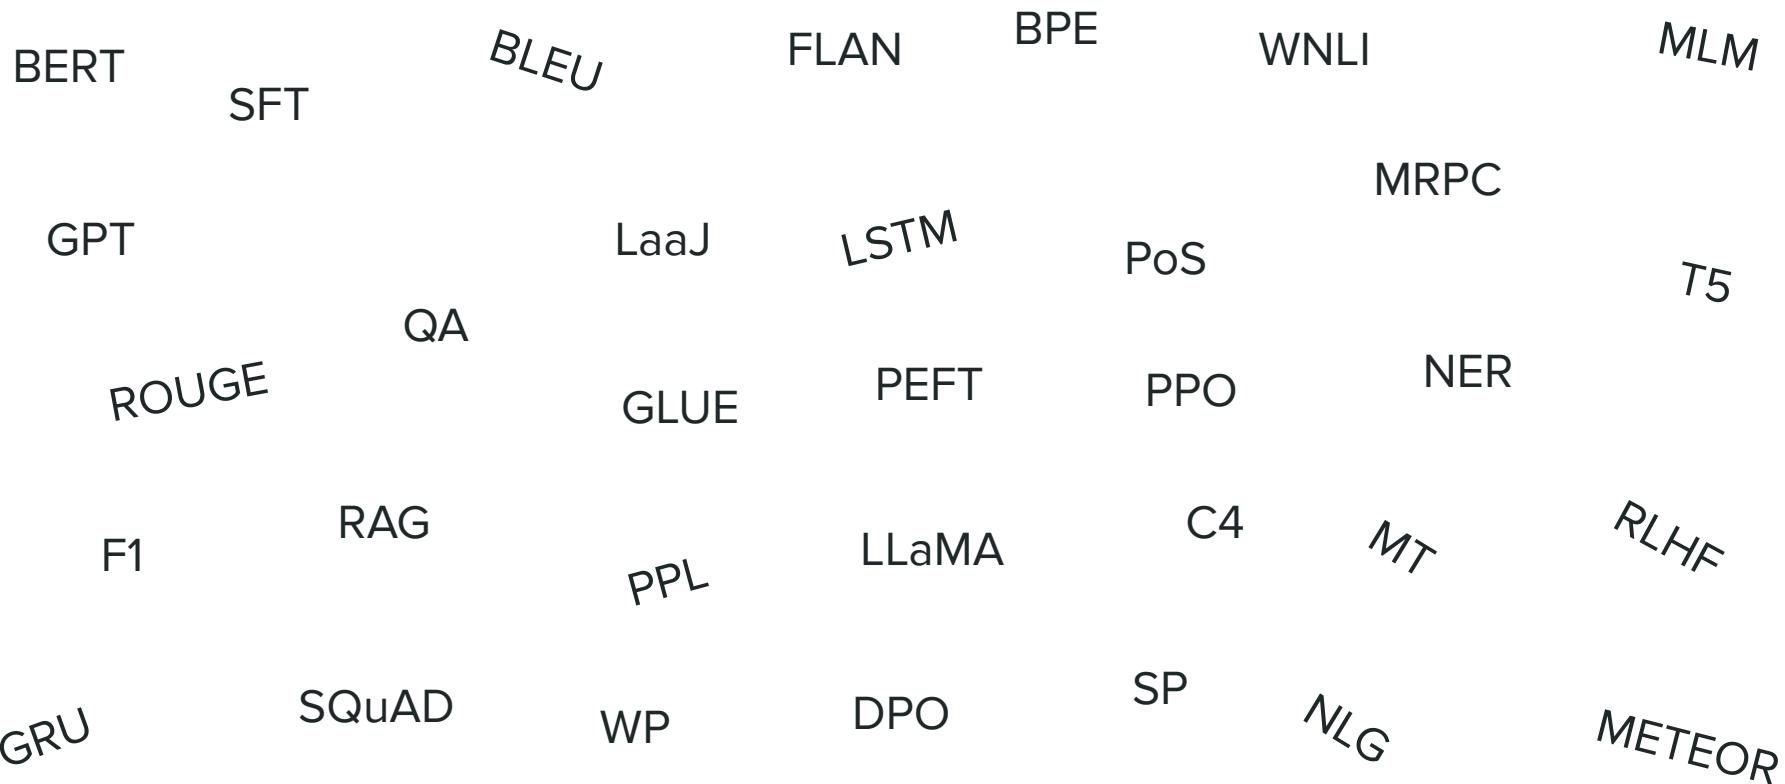
Example of a slide in this class

Explanation of a CME 295 concept

Source & suggested reading, if interested



Disclaimer before starting: many abbreviations....



...but don't worry!

BERT, T5, GPT, LLaMA

Transformer-based models

SFT, PEFT, FLAN, RL, RM, RLHF, PPO, DPO

Training strategies

MNLI, WNLI, C4, SQuAD, GLUE, MRPC

Datasets

LSTM, GRU, GloVe, BPE, CoT, ToT, SC, RAG

Misc architectures & techniques

NER, PoS, MLM, NSP, MT, QA, NLG

Tasks

F1, PPL, ROUGE, BLEU, METEOR, LaaJ, WER

Metrics



CME 295

Transformers & Large Language Models

NLP overview

Tokenization

Word representation

RNNs

Self-attention mechanism

Transformer architecture

End-to-end example

NLP tasks overview

Classification



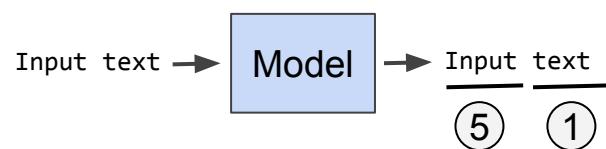
- Sentiment extraction
- Intent detection
- Language detection
- Topic modeling

NLP tasks overview

Classification



“Multi”-classification



- Sentiment extraction
- Intent detection
- Language detection
- Topic modeling

- Part of speech tagging
- **Named entity recognition**
- Dependency parsing
- Constituency parsing

linguistics side. not trendy

NLP tasks overview

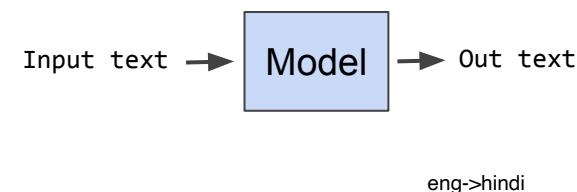
Classification



“Multi”-classification



Generation



- Sentiment extraction
- Intent detection
- Language detection
- Topic modeling

- Part of speech tagging
- Named entity recognition
- Dependency parsing
- Constituency parsing

- **Machine translation**
- **Question answering**
chatgpt, gmini
- **Summarization**
- **Text generation**
code gen

NLP task: Sentiment Extraction



NLP task: Sentiment Extraction



Datasets

 Amazon reviews

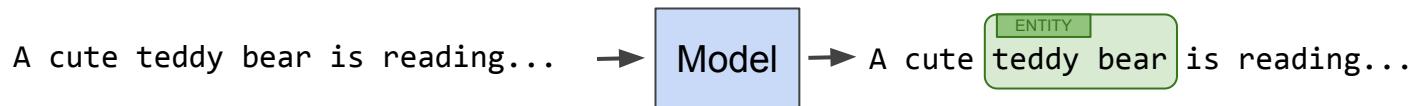
 IMDB critiques

 Twitter

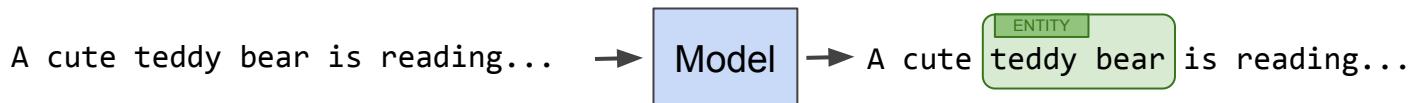
Evaluation metrics

- Accuracy → % of observations that were correctly predicted?
- Precision → % of predicted positive that were correct?
- Recall → % of actually positive that were correct?
- F1 score → score that is a function of precision and recall

NLP task: Named Entity Recognition



NLP task: Named Entity Recognition



Datasets



Annotated Reuters newspaper (CoNLL-2003, CoNLL++)

Evaluation metrics

- Accuracy
- Precision at a token level, per entity type
- Recall
- F1 score

NLP task: Machine Translation



NLP task: Machine Translation



Datasets

 WMT'14 English-French

 WMT'14 English-German

Evaluation metrics

bilingual evaluation under study

- BLEU → quality of text translated, similar to “precision”
need reference text or labels
- ROUGE → quality of text generated, similar to “recall”
- Perplexity → quantifies how ‘surprised’ the model is to see some words together
looks at the output probabilities

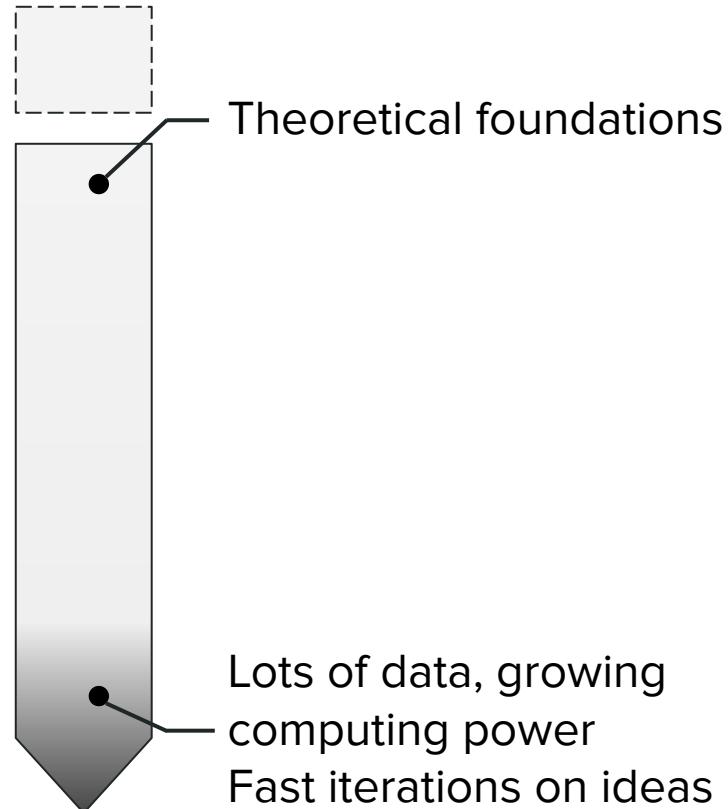
High-level timeline

1980s Recurrent neural networks (RNNs)

1997 Long short-term memory (LSTM)

2013
2017
2020s

meaningful embeddings
Word2vec
Transformers
Large Language Models





CME 295

Transformers & Large Language Models

NLP overview

Tokenization

Word representation

RNNs

Self-attention mechanism

Transformer architecture

End-to-end example

Tokenization

A cute teddy bear is reading.

Tokenization

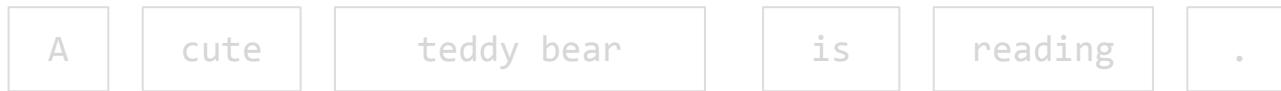
A cute teddy bear is reading.

arbitrary

A cute teddy bear is reading .

Tokenization

A cute teddy bear is reading.



Tokenization

A cute teddy bear is reading.

A cute teddy bear is reading .

A cute teddy bear is reading .

sub-word

A cute ted ##dy bear is read ##ing .

Tokenization

A cute teddy bear is reading.

A cute teddy bear is reading .

A cute teddy bear is reading .

A cute ted ##dy bear is read ##ing .

A - c u t e - t e d d y - b e a r - i s - r e a d i n g .

Tokenization summary

Method	Pros	Cons
Word-level	<ul style="list-style-type: none">SimpleInterpretable	<ul style="list-style-type: none">Risk of OOV <small>out of vocabulary</small>Does not leverage knowledge of root
Subword-level e.g. WordPiece, BPE	<ul style="list-style-type: none">Leverages common prefixes and suffixesLearned from the data	<ul style="list-style-type: none">Risk of OOV, though less than word-level
Character-level	<ul style="list-style-type: none">Small chance of OOVRobust to Capitalizing and Mispellings	<ul style="list-style-type: none">Makes computations slowerEmbeddings not interpretable



CME 295

Transformers & Large Language Models

NLP overview

Tokenization

Word representation

RNNs

Self-attention mechanism

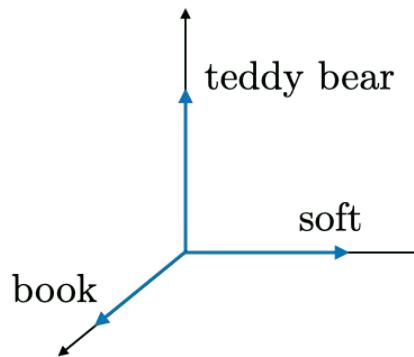
Transformer architecture

End-to-end example

Token representations

Motivation

Naive (one-hot) encoding

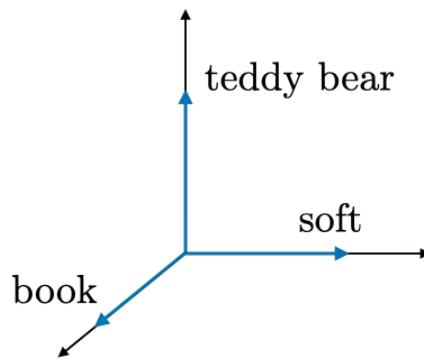


$$\text{soft} = \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix} \quad \langle \text{teddy bear}, \text{ book} \rangle = 0 \\ \langle \text{teddy bear}, \text{ soft} \rangle = 0$$

Token representations

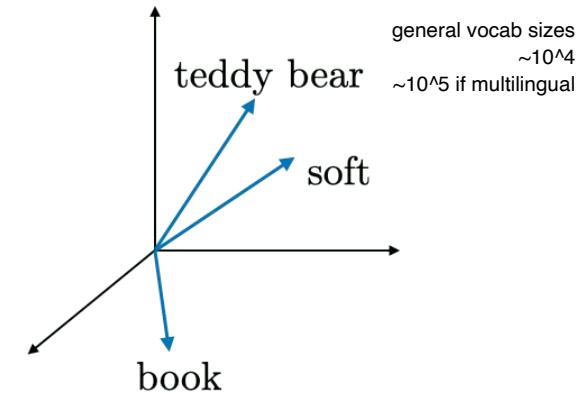
Motivation

Naive (one-hot) encoding



$$\text{soft} = \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix} \quad \langle \text{teddy bear}, \text{ book} \rangle = 0 \\ \langle \text{teddy bear}, \text{ soft} \rangle = 0$$

Learned embedding



$$\text{soft} = \begin{pmatrix} 0.95 \\ 0.32 \\ 0.01 \end{pmatrix} \quad \langle \text{teddy bear}, \text{ book} \rangle \sim 0 \\ \langle \text{teddy bear}, \text{ soft} \rangle \sim 1$$

general vocab sizes
~ 10^4
~ 10^5 if multilingual

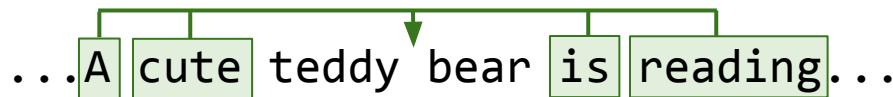
Word2vec

Overview

- Neural network with a **proxy task** over billions of words worth of text
- Learns an embedding layer

Proxy tasks

- CBOW (continuous bag of words)



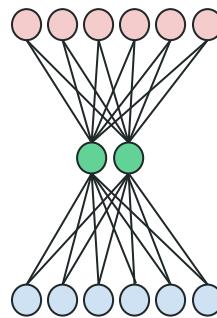
- Skip-gram



Word2vec

Architecture

output



size V

hidden

size d

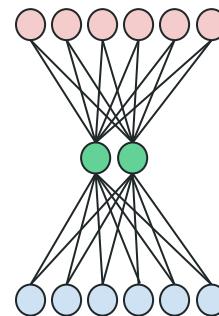
input

size V

Word2vec

Example with predicting next word

A **cute** teddy bear is reading

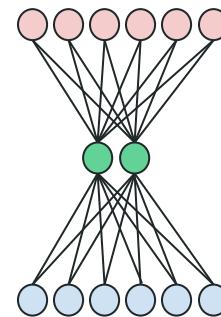


A **cute** teddy bear is reading

Word2vec

Example with predicting next word

A **cute** teddy bear is reading



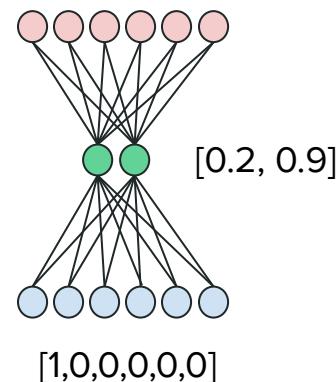
[1,0,0,0,0,0]

A **cute** teddy bear is reading

Word2vec

Example with predicting next word

A cute teddy bear is reading



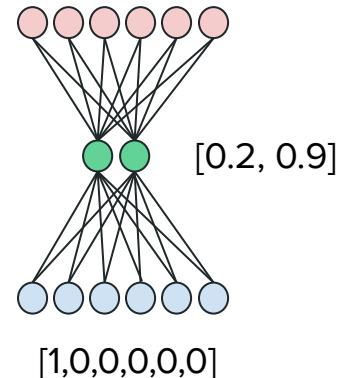
A cute teddy bear is reading

Word2vec

Example with predicting next word

A cute teddy bear is reading

[0.2, 0.4, 0.1, 0.1, 0.1, 0.1]

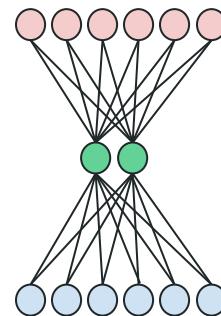


A cute teddy bear is reading

Word2vec

Example with predicting next word

A cute **teddy bear** is reading

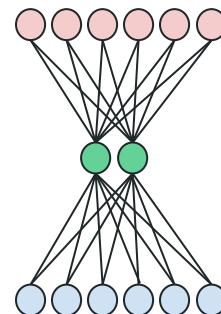


A **cute** teddy bear is reading

Word2vec

Example with predicting next word

A cute **teddy bear** is reading



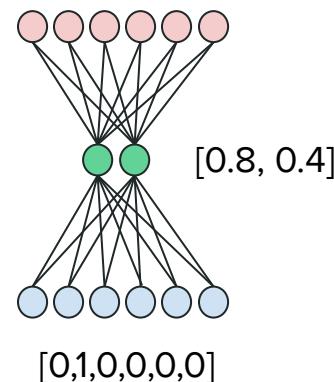
[0,1,0,0,0,0]

A **cute** teddy bear is reading

Word2vec

Example with predicting next word

A cute **teddy bear** is reading



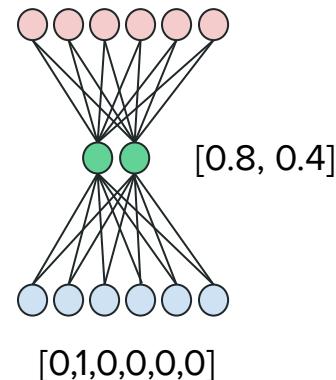
A **cute** teddy bear is reading

Word2vec

Example with predicting next word

A cute **teddy bear** is reading

[0.2, 0.2, 0.2, 0.1, **0.2**, 0.1]

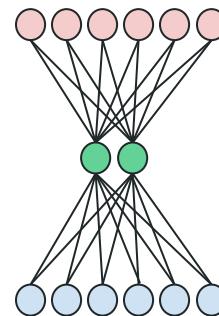


A **cute** teddy bear is reading

Word2vec

Example with predicting next word

A cute teddy bear **is** reading



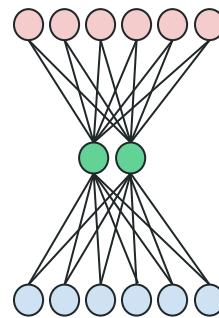
A cute **teddy bear** is reading

Word2vec

Example with predicting next word

A cute teddy bear is **reading**

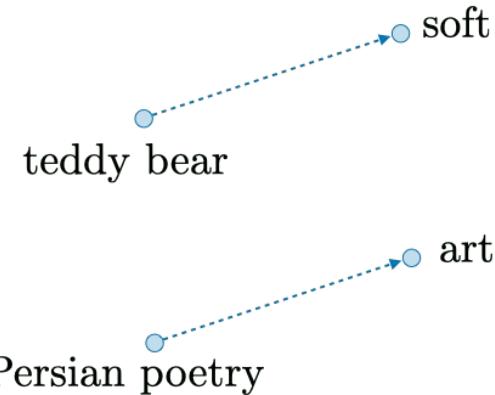
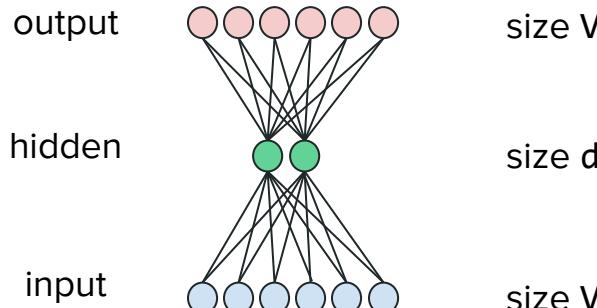
length of embedding is 10^2 - 10^3 now a days. 768 empirical. and there is a trade-off in how long you want it



A cute teddy bear **is** reading

Word2vec

A **cute** teddy bear is reading



A **cute** teddy bear is reading

representations of sentences or pieces of texts. naive-take avg.
loose meaning.

representations are token specific.
we need to capture the sequential nature of text



CME 295

Transformers & Large Language Models

NLP overview

Tokenization

Word representation

RNNs

Self-attention mechanism

Transformer architecture

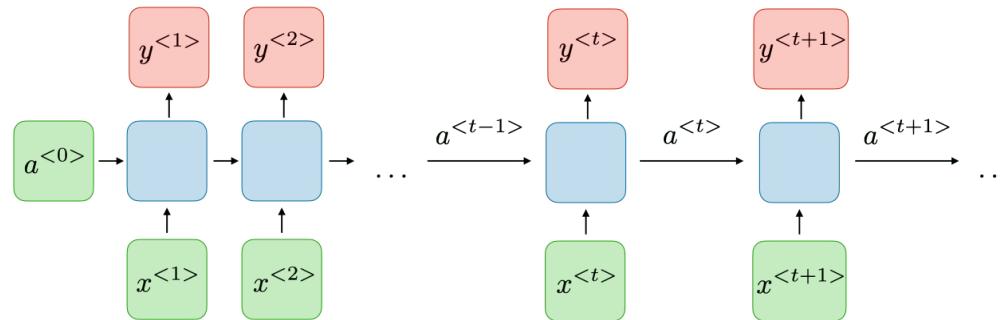
End-to-end example

Recurrent Neural Networks (RNNs)

Overview

- First introduced in the 80s
- Class of neural networks where connections form a temporal sequence

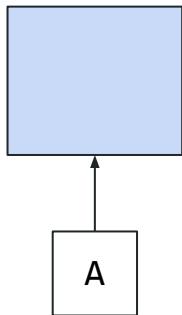
General form



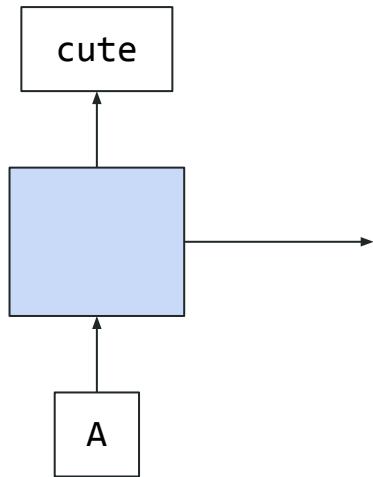
Recurrent Neural Networks (RNNs)

A

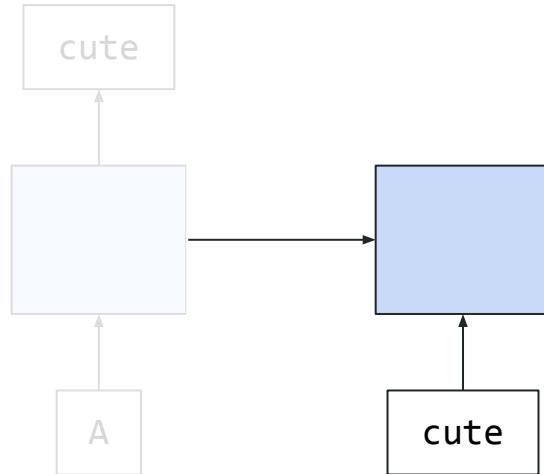
Recurrent Neural Networks (RNNs)



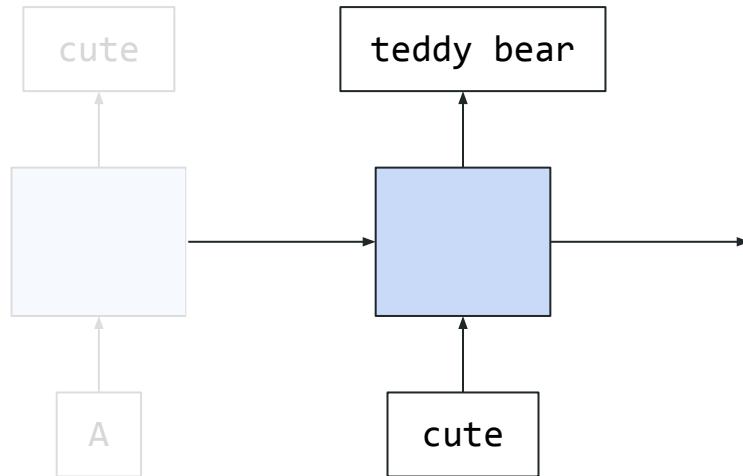
Recurrent Neural Networks (RNNs)



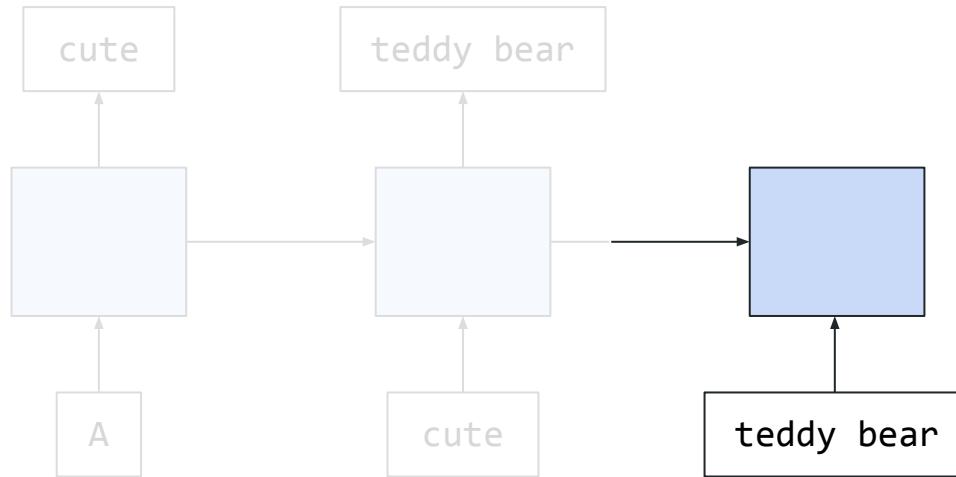
Recurrent Neural Networks (RNNs)



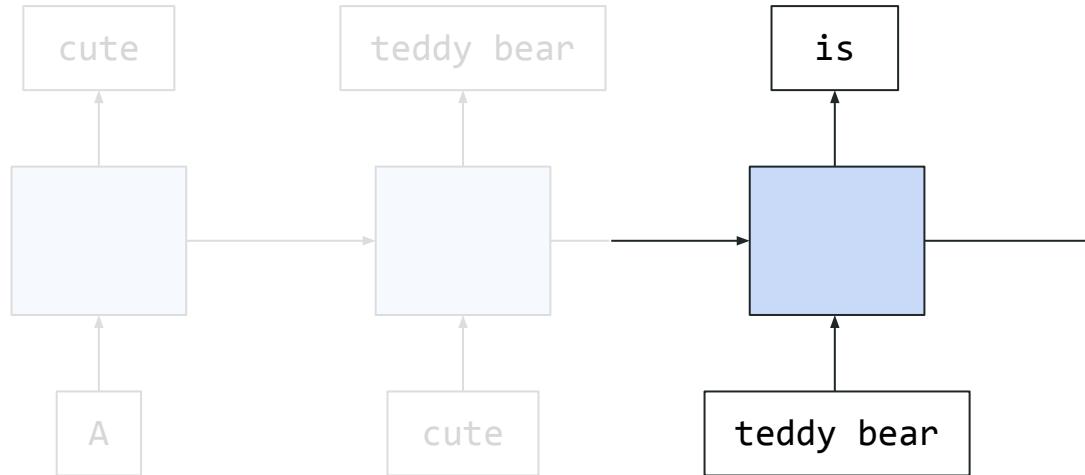
Recurrent Neural Networks (RNNs)



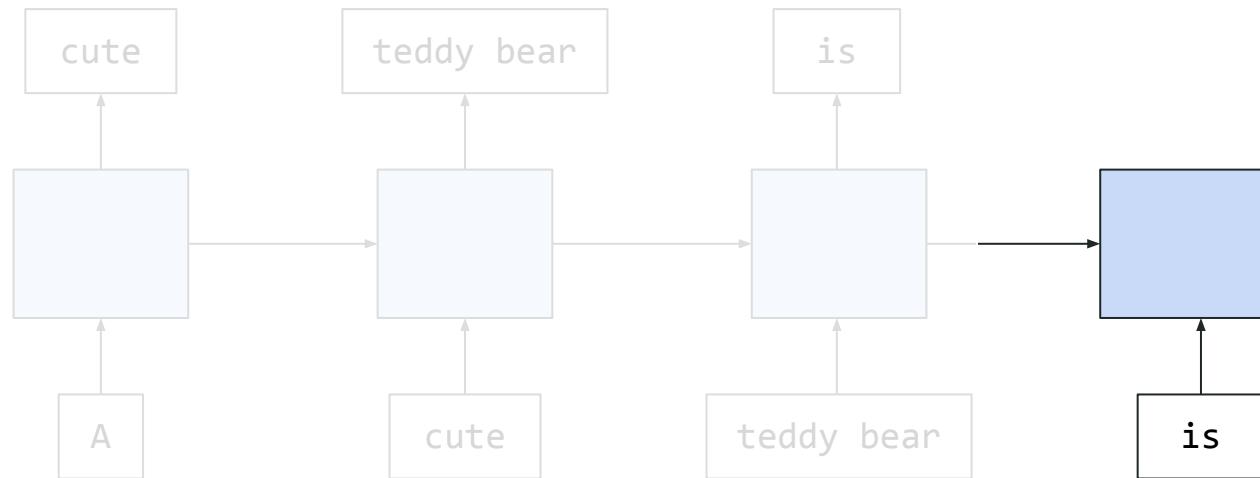
Recurrent Neural Networks (RNNs)



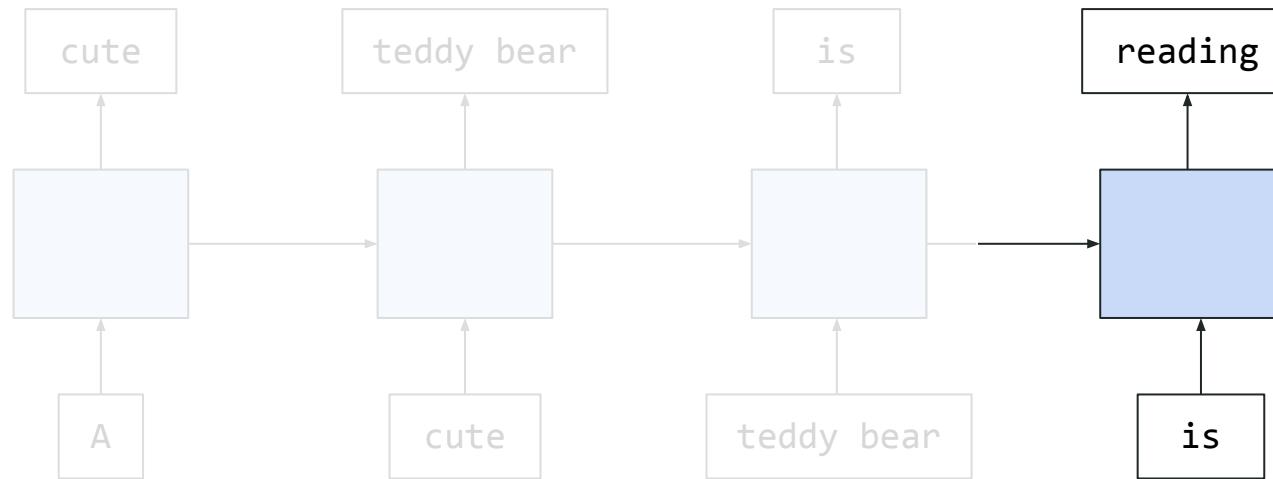
Recurrent Neural Networks (RNNs)



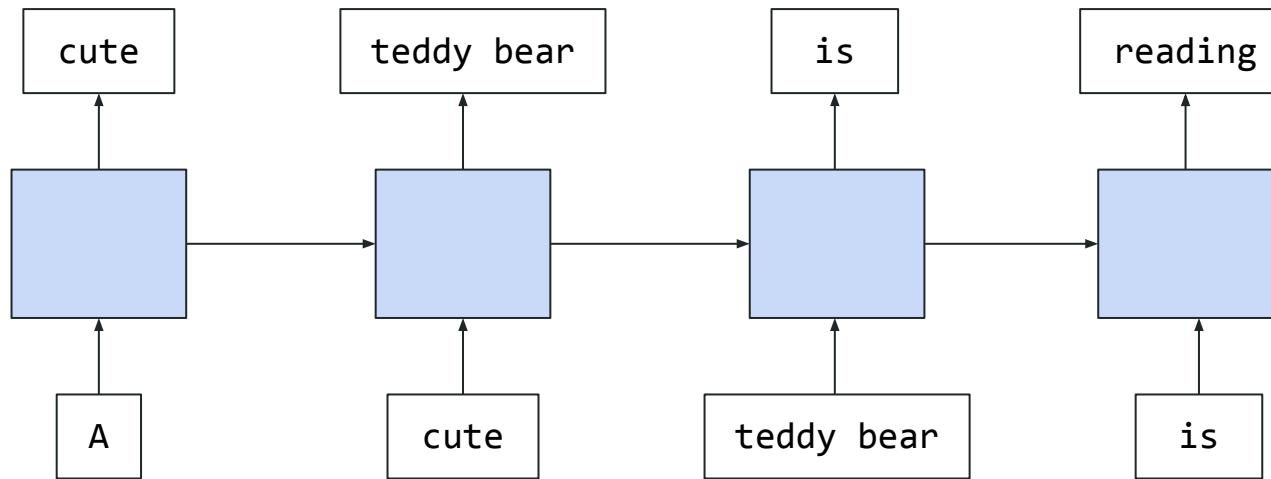
Recurrent Neural Networks (RNNs)



Recurrent Neural Networks (RNNs)



Recurrent Neural Networks (RNNs)

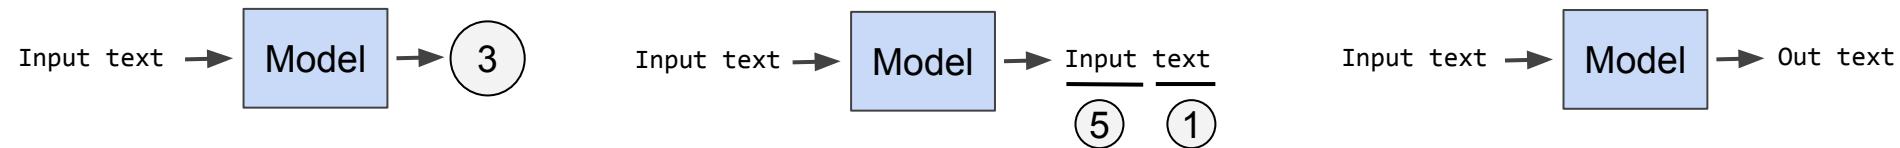


Recurrent Neural Networks (RNNs)

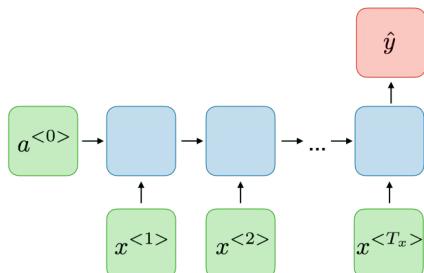
Classification

“Multi”-classification

Generation

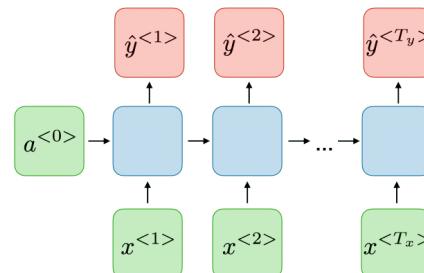


Sentiment



Opinion

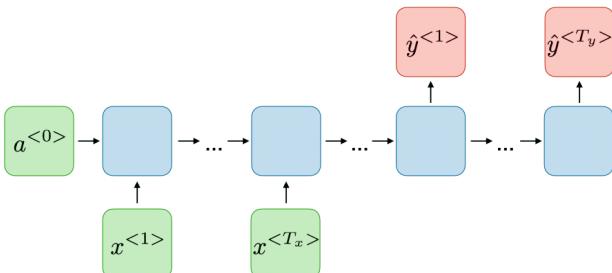
Tags



Text

meaning is solely in the hidden states.
long range dependancies. impacts ability to remember
what model saw in the past
vanishing gradient

Translation



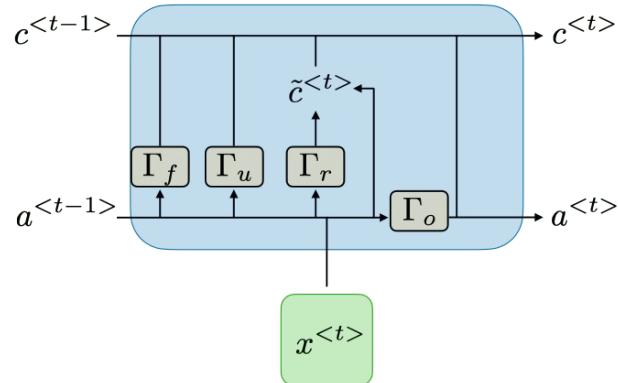
Source

Long Short-Term Memory (LSTM)

Overview

- Introduced in “[Long short-term memory](#)” (1997)
- Uses a more structured approach in the cell’s hidden state

General form



Summary of main methods (non-exhaustive list)

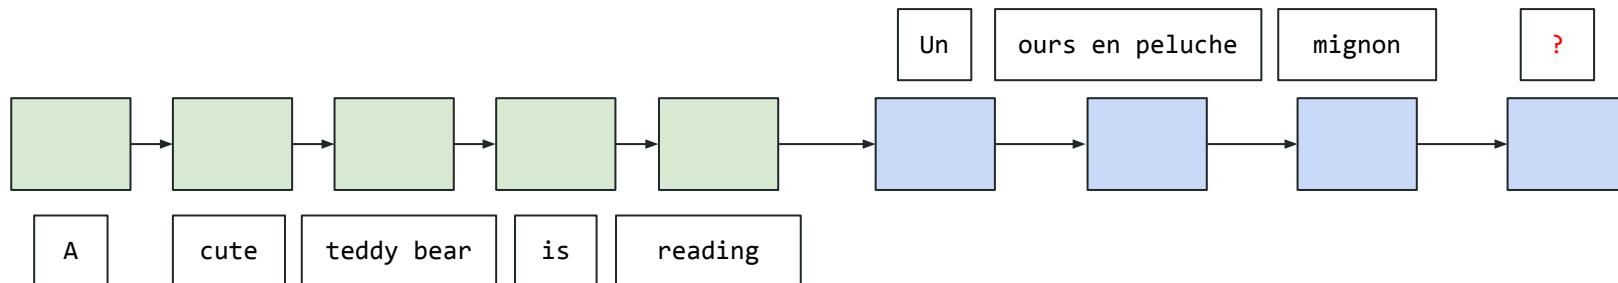
Method	Pros	Cons
Word2vec e.g. CBOW, Skip-gram	<ul style="list-style-type: none">• Very simple, yet powerful• Intuitive embeddings	<ul style="list-style-type: none">• Word order does not count• Embeddings not context aware
Recurrent Neural Networks e.g. traditional RNN, LSTM	<ul style="list-style-type: none">• Word order matters• State-of-the-art results	<ul style="list-style-type: none">• Vanishing gradient problem• Slow computations

History of attention

- Introduced in 2014
- Translation tasks had a real issue with long-term dependencies
- Seq2seq unable to "remember" what input sentence was saying

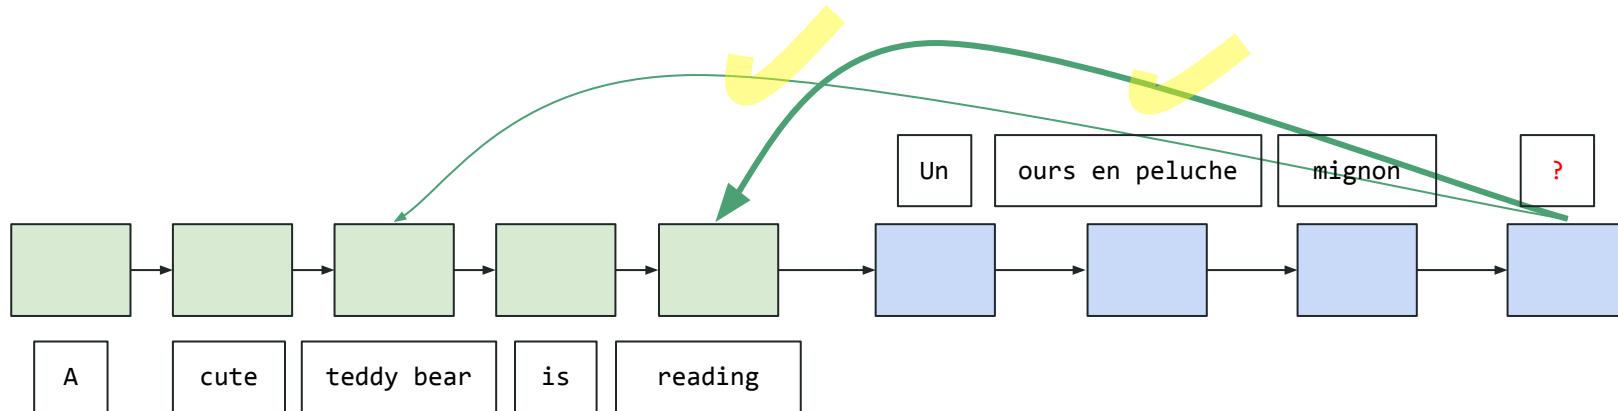
History of attention

- Introduced in 2014
- Translation tasks had a real issue with long-term dependencies
- Seq2seq unable to "remember" what input sentence was saying



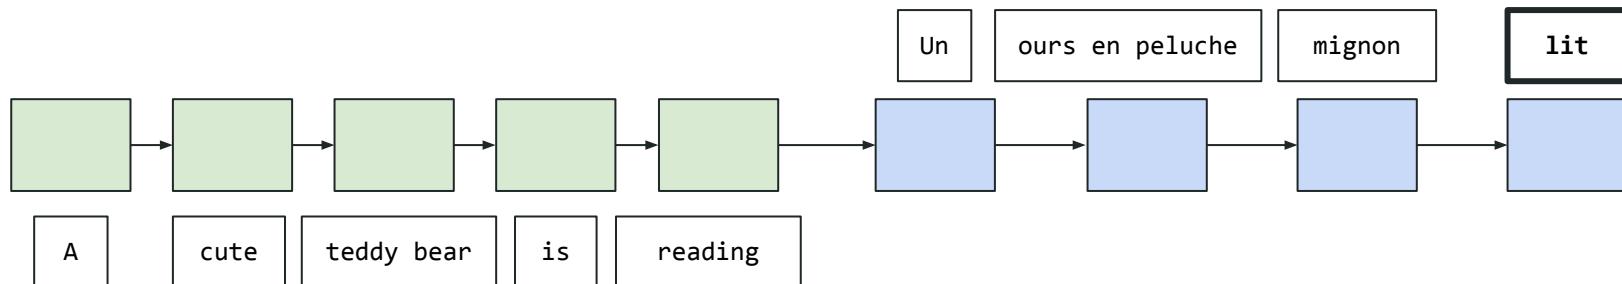
History of attention

- Introduced in 2014
- Translation tasks had a real issue with long-term dependencies
- Seq2seq unable to "remember" what input sentence was saying



History of attention

- Introduced in 2014
- Translation tasks had a real issue with long-term dependencies
- Seq2seq unable to "remember" what input sentence was saying





Transformers & Large Language Models

NLP overview

Tokenization

Word representation

RNNs

Self-attention mechanism

Transformer architecture

End-to-end example

Overview of the Transformer

- Introduced in the 2017 paper "**Attention is All You Need**"
- Relies on the **self-attention** mechanism
- State-of-the-art results on machine translation tasks

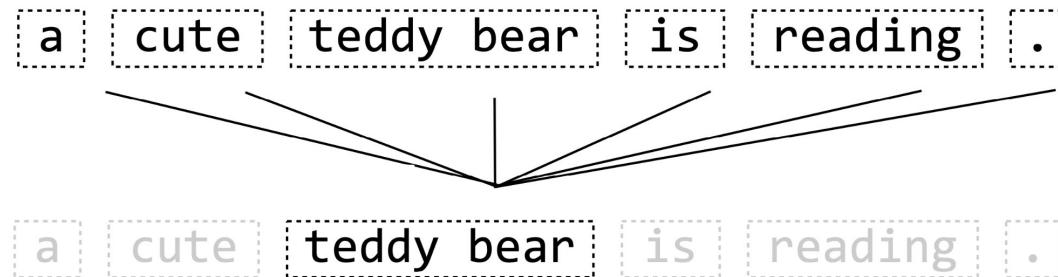
Overview of the Transformer

- Introduced in the 2017 paper "**Attention is All You Need**"
- Relies on the **self-attention** mechanism
- State-of-the-art results on machine translation tasks

a cute teddy bear is reading .

Overview of the Transformer

- Introduced in the 2017 paper "**Attention is All You Need**"
- Relies on the **self-attention** mechanism
- State-of-the-art results on machine translation tasks



Attention mechanism

Concept of **Query**, **Key**, **Value**

a cute teddy bear is reading .

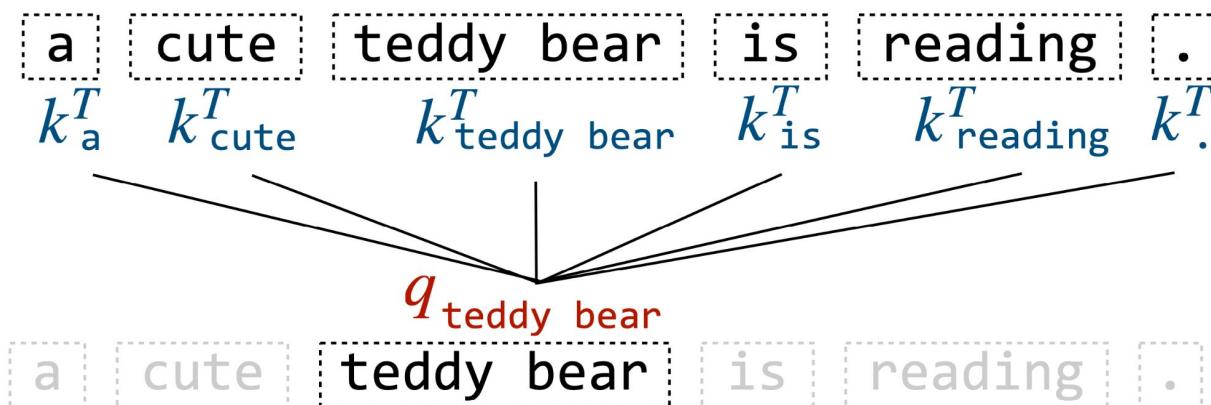
Attention mechanism

Concept of **Query**, **Key**, **Value**

$q_{\text{teddy bear}}$
a cute teddy bear is reading .

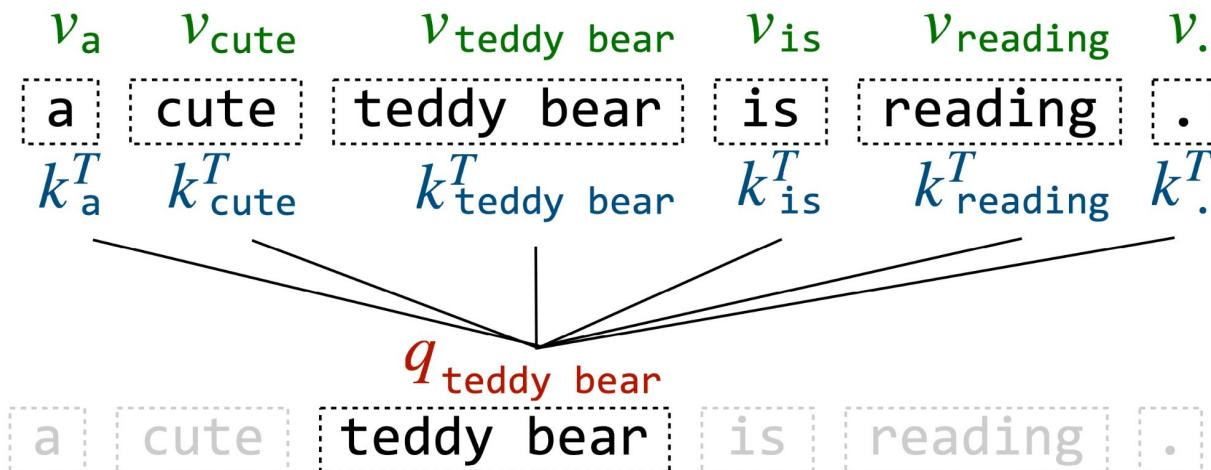
Attention mechanism

Concept of **Query**, **Key**, **Value**



Attention mechanism

Concept of **Query**, **Key**, **Value**



Attention mechanism

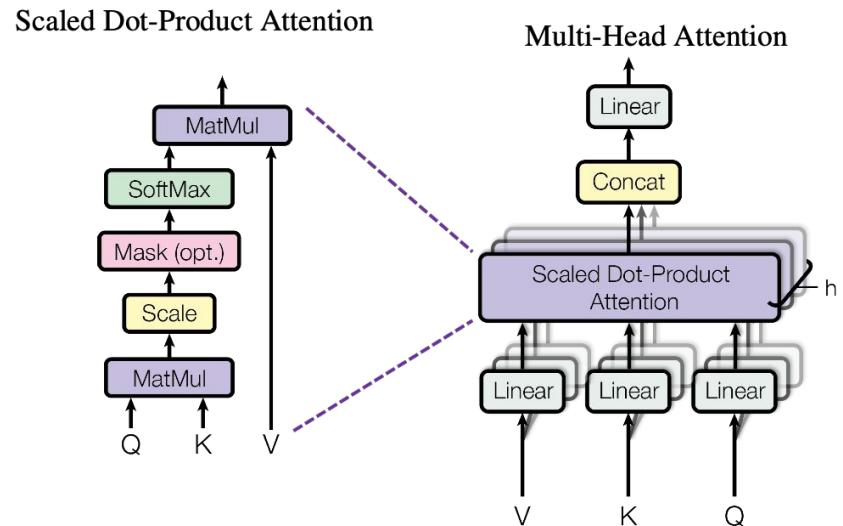
Efficient computations with matrices:

$$\text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

Attention mechanism

Efficient computations with matrices:

$$\text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$





Transformers & Large Language Models

NLP overview

Tokenization

Word representation

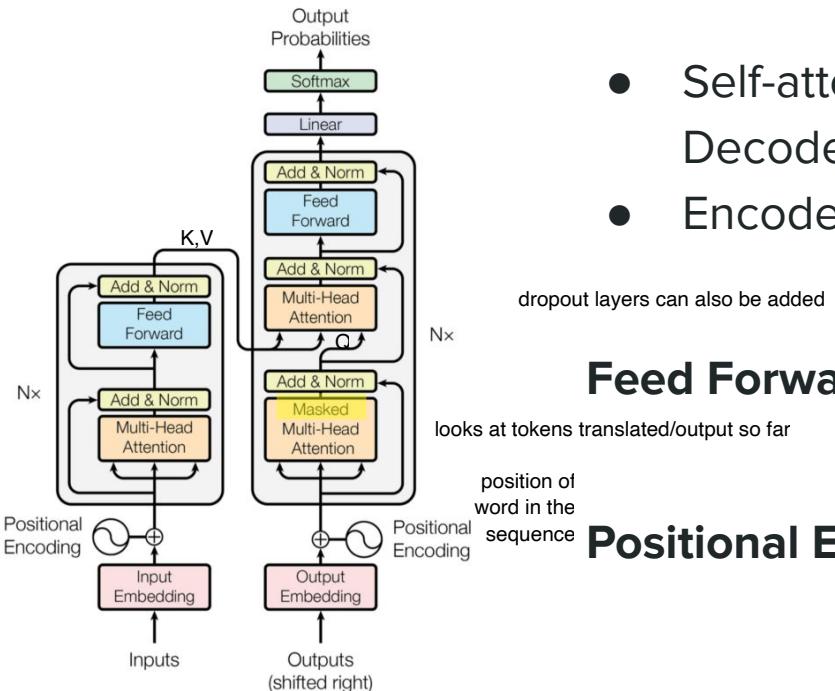
RNNs

Self-attention mechanism

Transformer architecture

End-to-end example

Transformer architecture



Attention layer (MHA)

- Self-attention (Encoder-Encoder, Decoder-Decoder)
- Encoder-Decoder attention layer

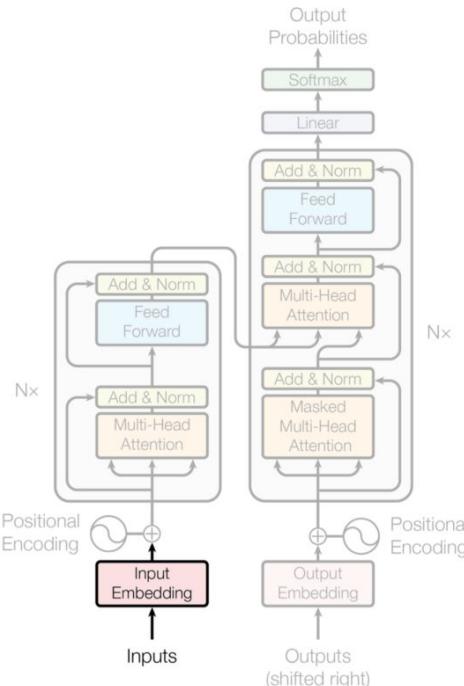
Feed Forward Neural Network (FFNN)

Positional Encoding (PE)

looks at tokens translated/output so far

position of
word in the
sequence

Input



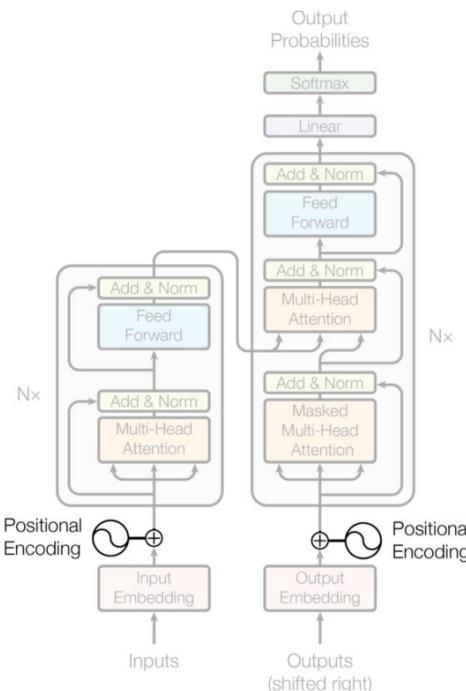
Overview

- Text is "tokenized"
- Learned embeddings for tokens

Parameters

- V : vocabulary size
- d_{model} : embedding dimensions

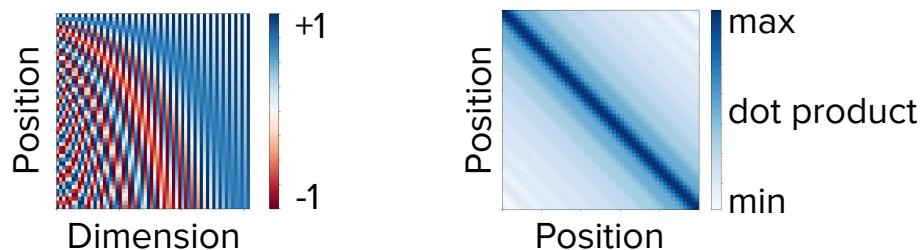
... with a trick!



Positional encoding

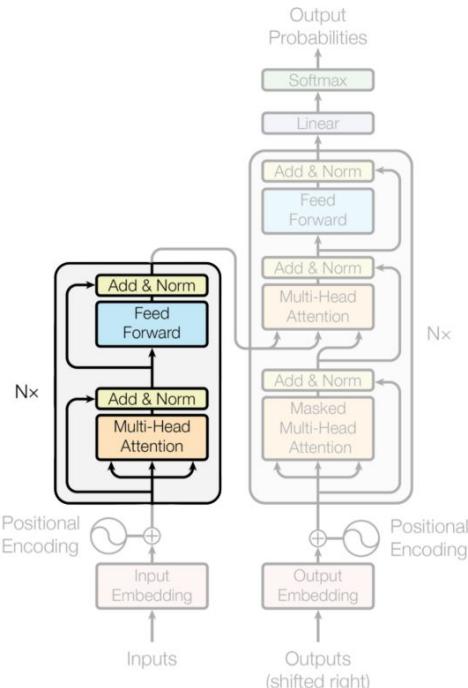
Idea:

- Add **position information** to inputs
- Can be either learned or hardcoded



Goal: let model understand relative input position

Encoder



Overview

tries to figure out how to express things as a function
of other things in the input sequence

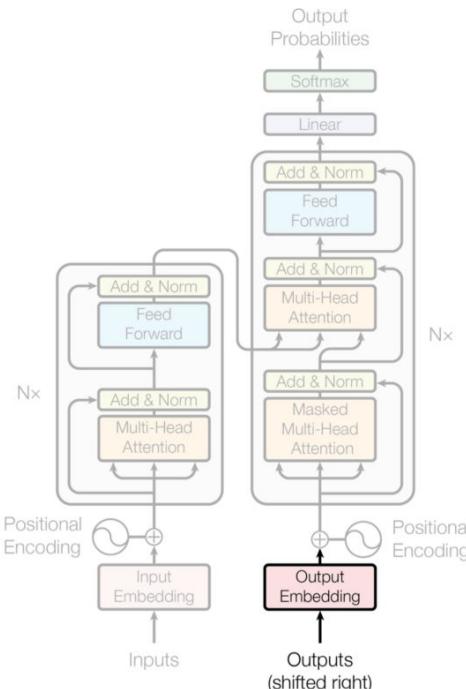
- Encoder-Encoder attention / self-attention
- Feed Forward Neural Network
- Normalization layer

project for higher degrees of freedom to learn

Parameters

- N : layers stacked
- h : number of attention heads
- d_{FF} , d_{key} , d_{value} : sub-layer dimension
- d_{model} : embedding dimensions

Output "shifted right"



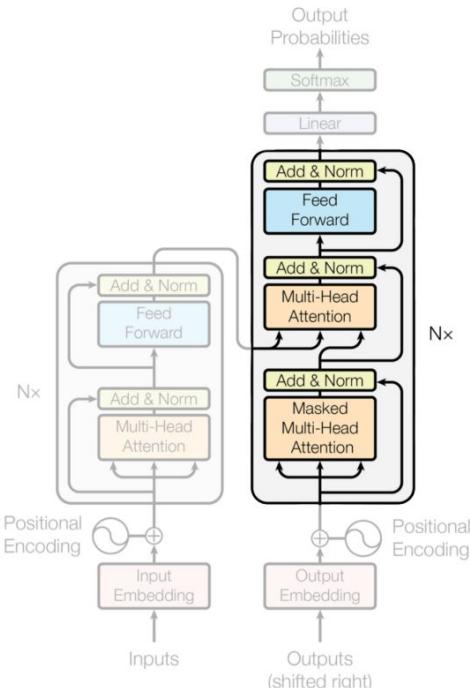
Overview

- Learned embeddings for output tokens
- In practice, will start with [BOS] during translation

Parameters

- V : vocabulary size
- d_{model} : embedding dimensions

Decoder



Overview

tries to figure out what is useful from what
has already been output

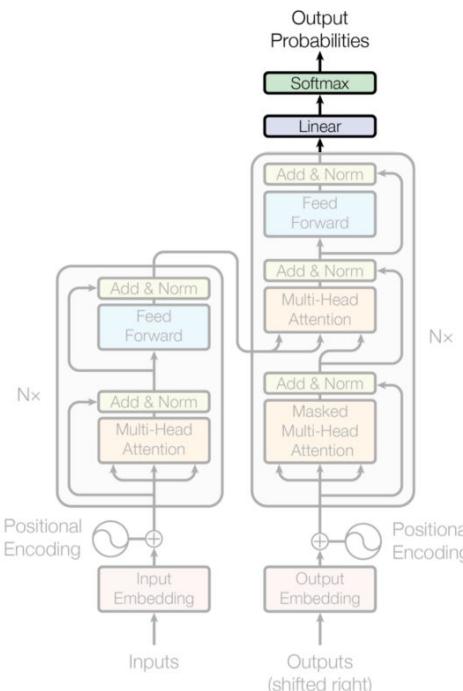
- Decoder-Decoder attention / self-attention
- Encoder-Decoder attention express things in a function of what was in input cross-attention layer
- Feed Forward Neural Network
- Normalization layer

Parameters

- N: layers stacked
- h: number of attention heads
- d_{FF} , d_{key} , d_{value} : sub-layer dimension
- d_{model} : embedding dimensions

Output

Overview

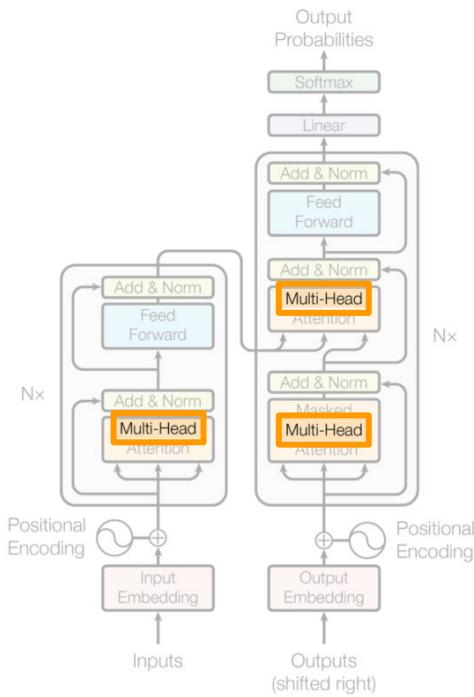


- Linear projection
- Classification problem that outputs probability of belonging to a class, where class = word

Parameters

- V : vocabulary size
- d_{model} : embedding dimensions

Computational tricks



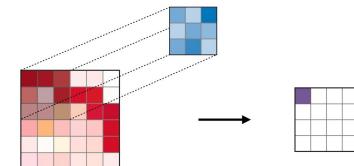
Multi-head attention

Idea:

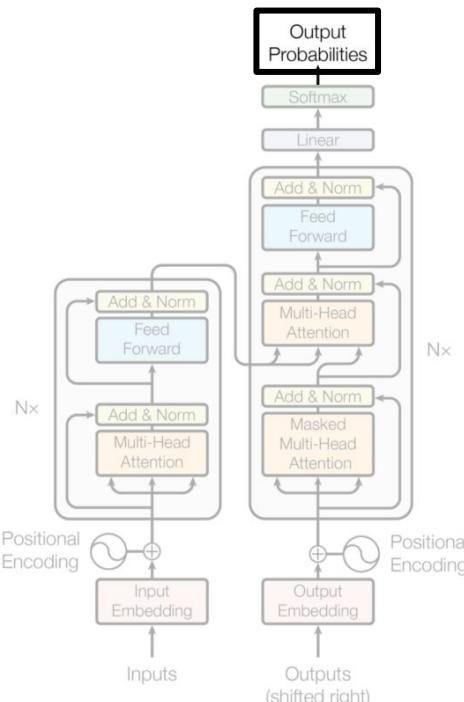
- Run **multiple** self-attention layers in parallel

Benefits:

- Enables the model to capture different attention features in parallel
- Comparison: **multiple** filters of a convolutional layer in computer vision



Computational tricks



Label smoothing

Idea:

instead of $(1,0,0,0)$ $(1-\epsilon, \epsilon_1, \epsilon_2, \epsilon_3)$

- **2015 vision paper**: overconfidence is bad
- Introduce **noise** in true labels

$$q(k|x) = \delta_{k,y} \rightarrow q'(k|x) = (1 - \epsilon)\delta_{k,y} + \epsilon u(k)$$

Benefits:

- General technique that prevents overfitting
- Improves accuracy and BLEU score



Transformers & Large Language Models

NLP overview

Tokenization

Word representation

RNNs

Self-attention mechanism

Transformer architecture

End-to-end example

Stitching all the pieces together with an example

A cute teddy bear is reading.

Stitching all the pieces together with an example

A cute teddy bear is reading .

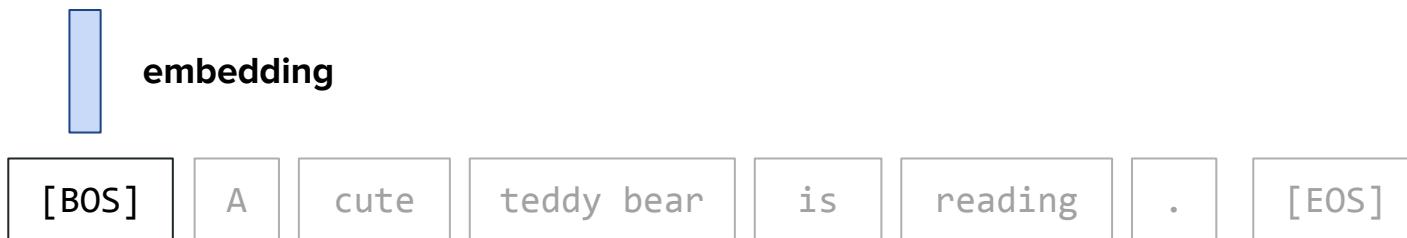
Stitching all the pieces together with an example

[BOS] A cute teddy bear is reading . [EOS]

Stitching all the pieces together with an example



Stitching all the pieces together with an example



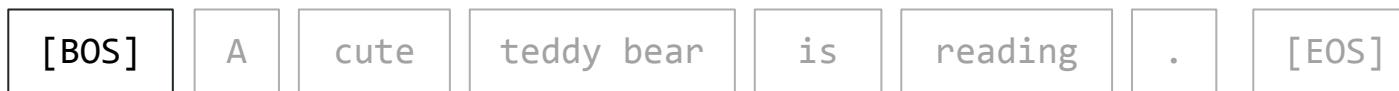
Stitching all the pieces together with an example



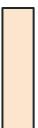
position embedding



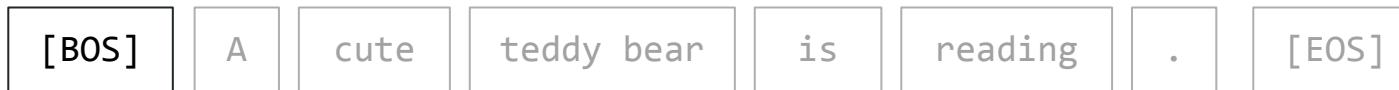
embedding



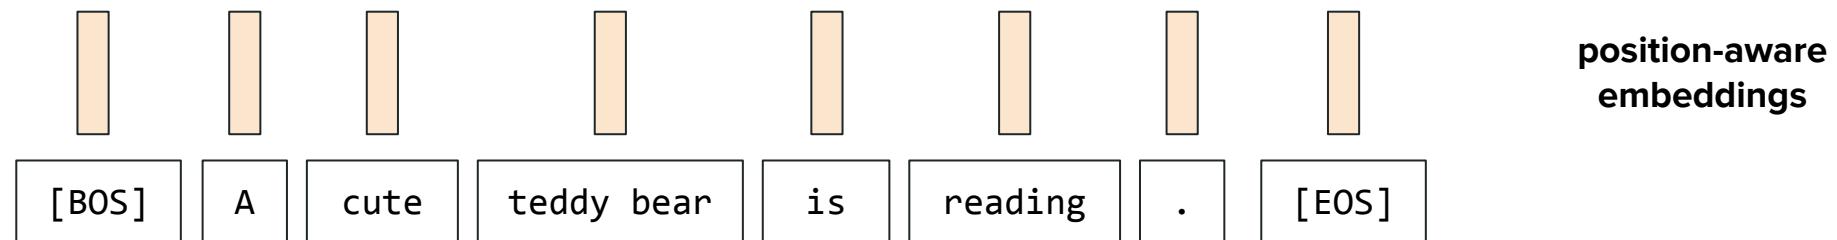
Stitching all the pieces together with an example



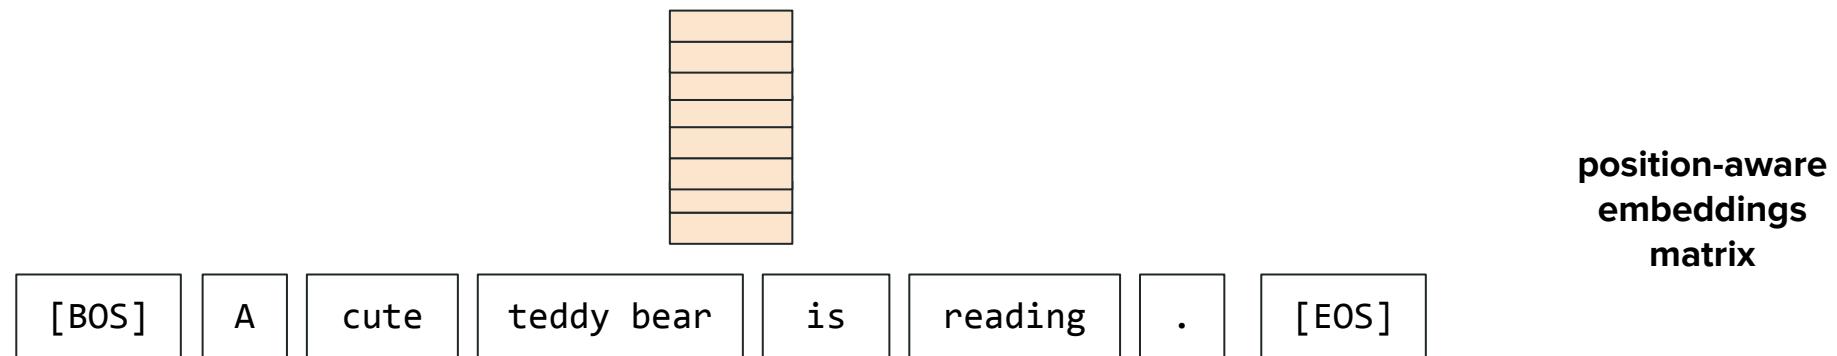
position-aware embedding



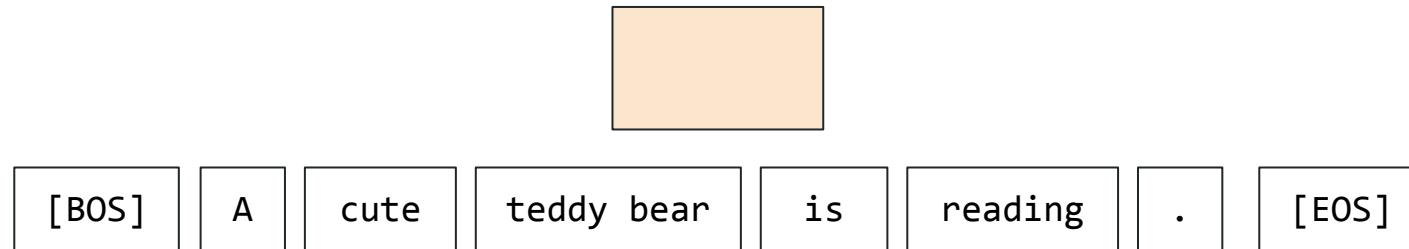
Stitching all the pieces together with an example



Stitching all the pieces together with an example



Stitching all the pieces together with an example



**position-aware
embeddings
matrix**

Stitching all the pieces together with an example



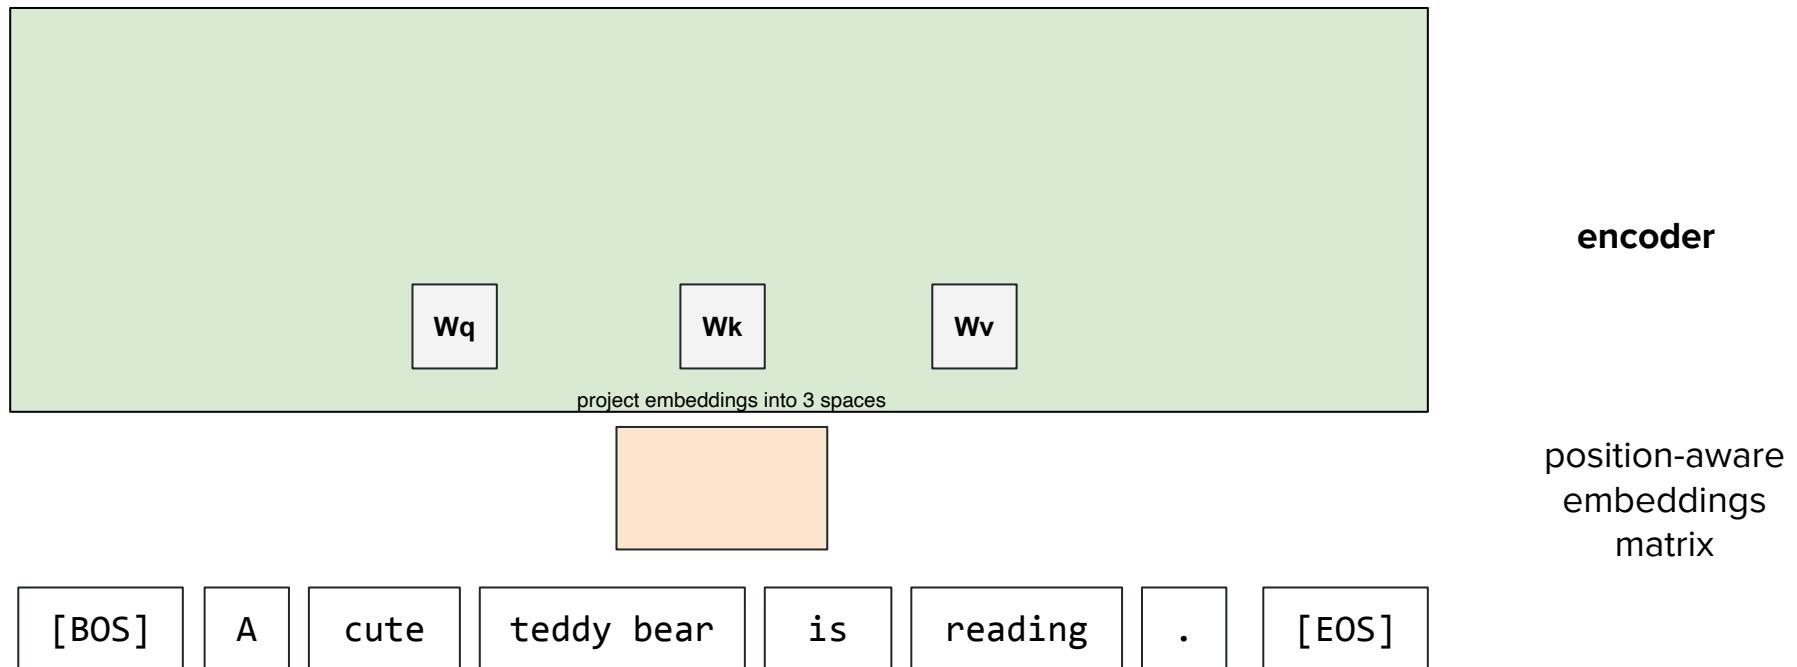
encoder



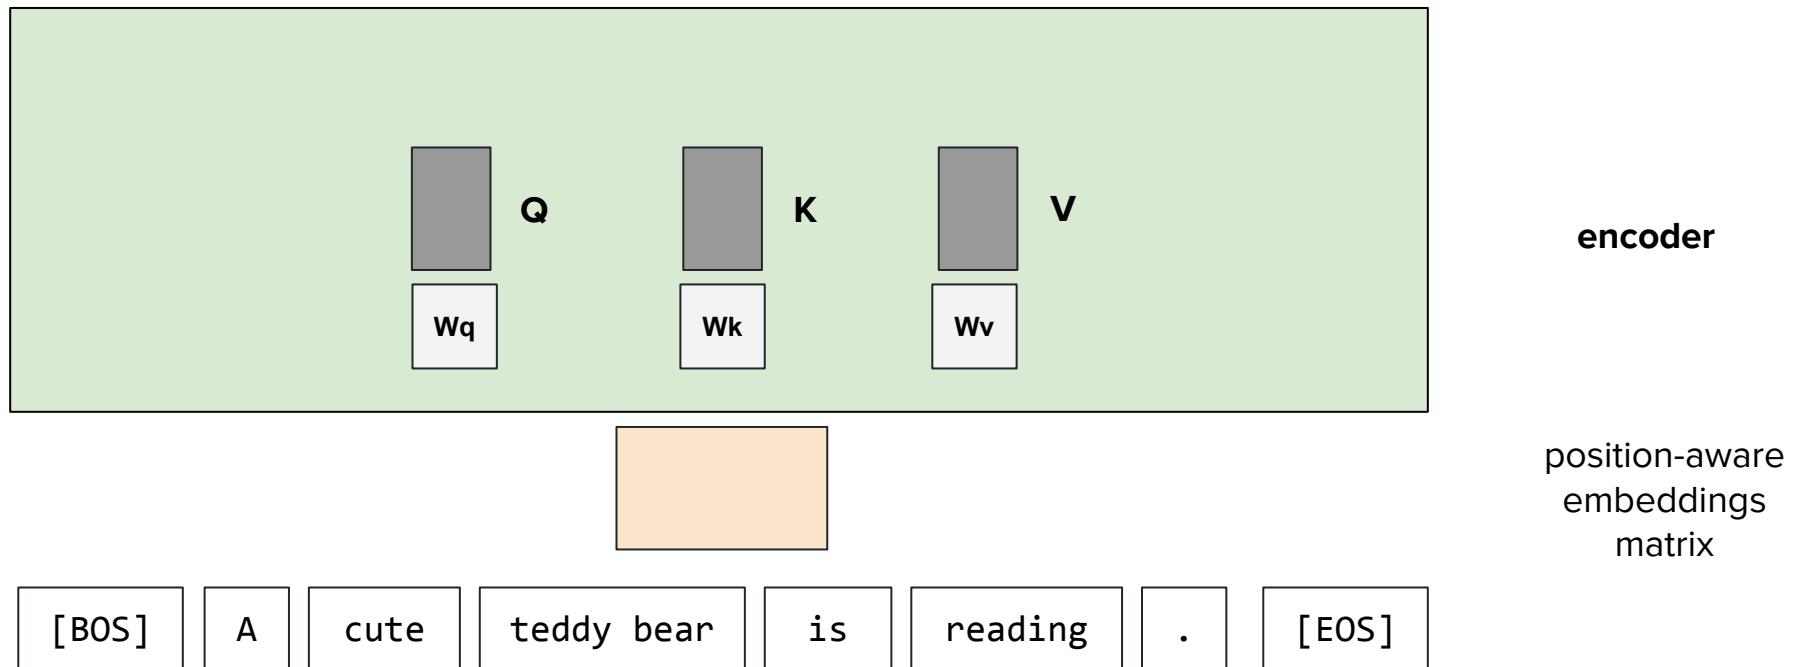
position-aware
embeddings matrix



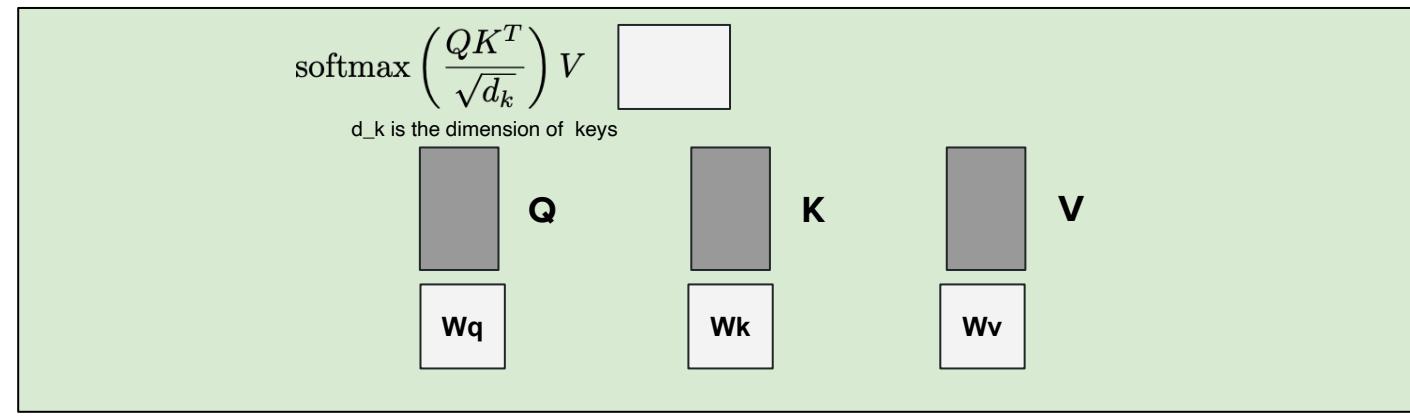
Stitching all the pieces together with an example



Stitching all the pieces together with an example



Stitching all the pieces together with an example

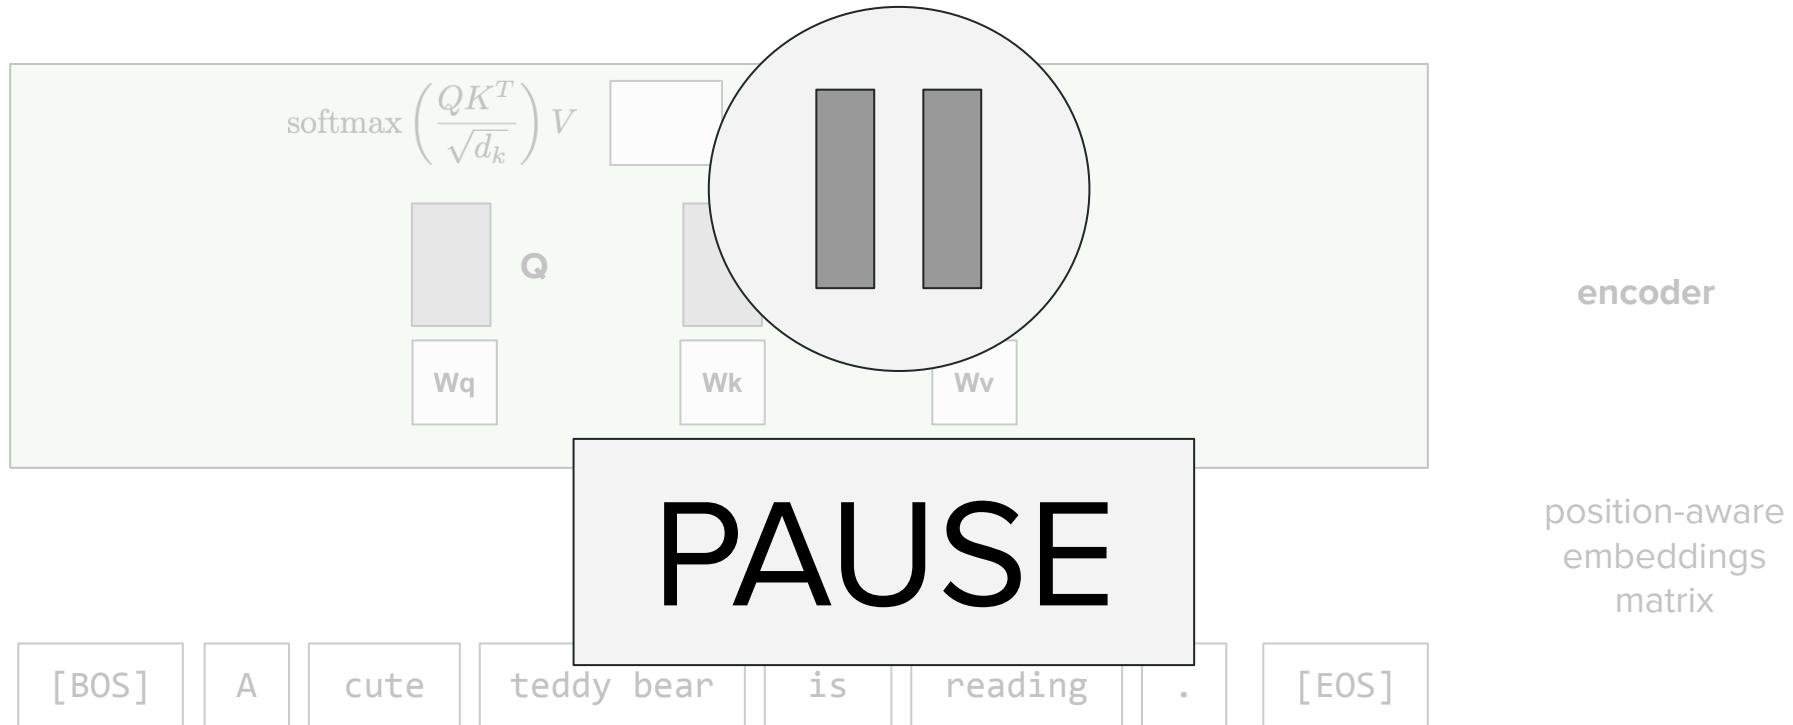


encoder

position-aware
embeddings
matrix

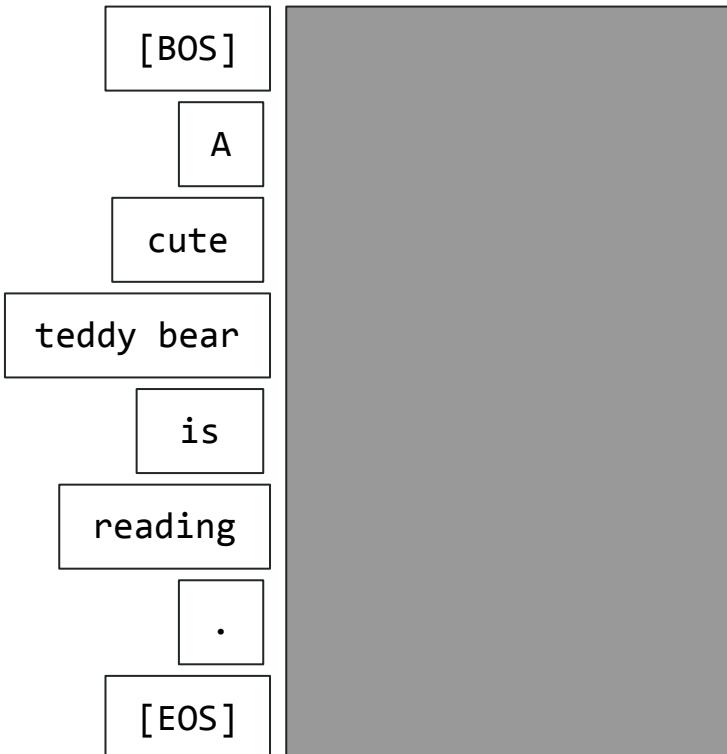


Stitching all the pieces together with an example

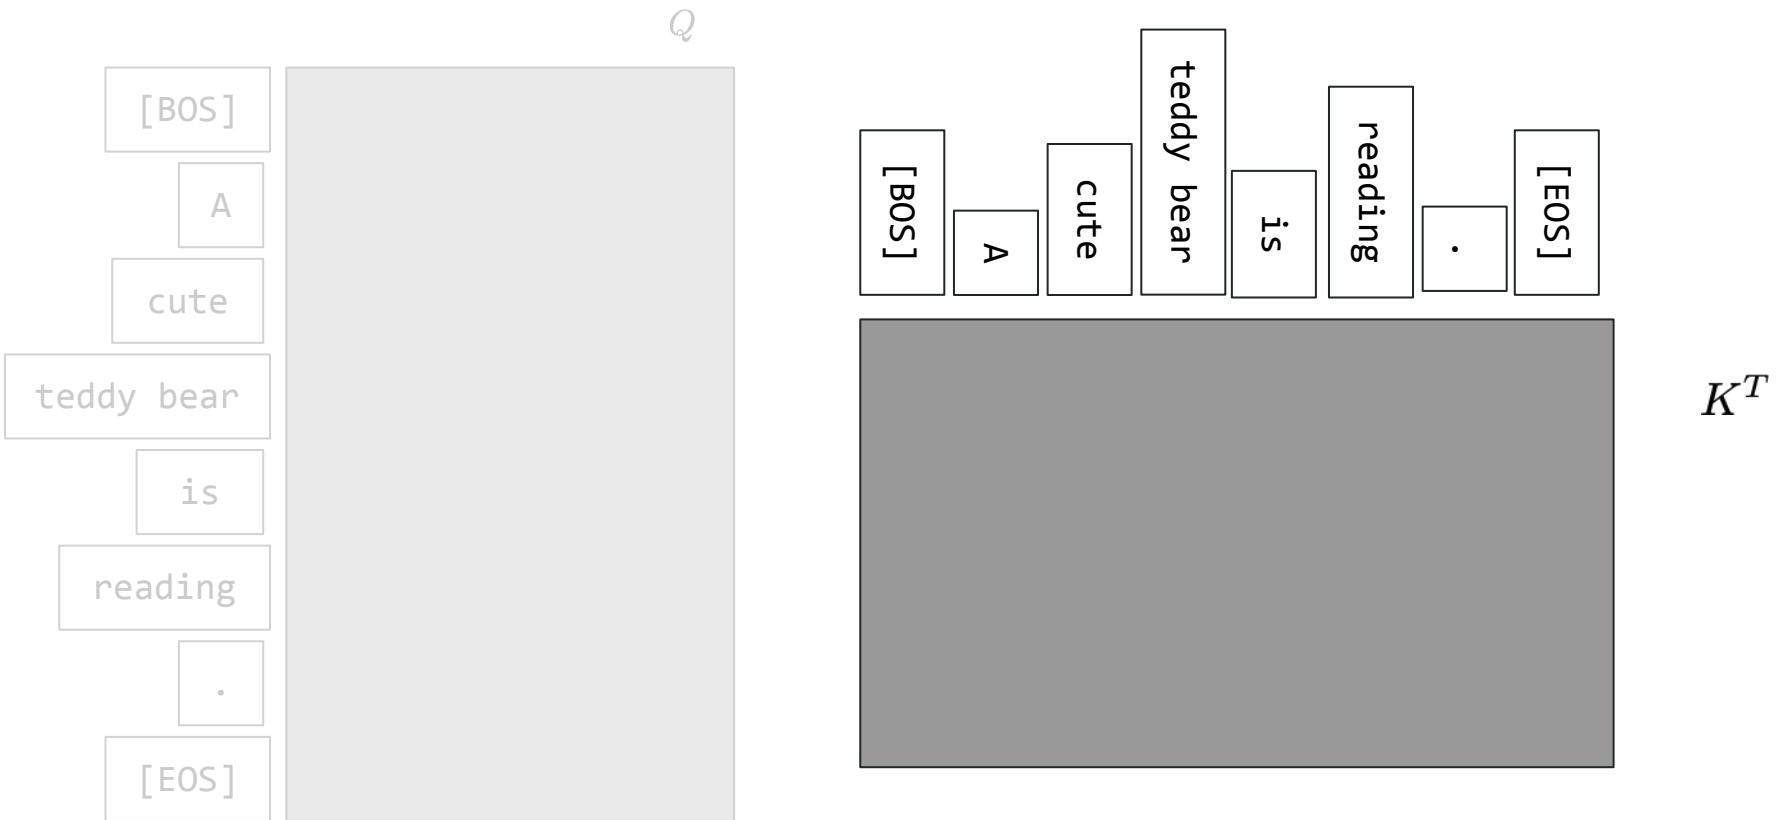


Stitching all the pieces together with an example

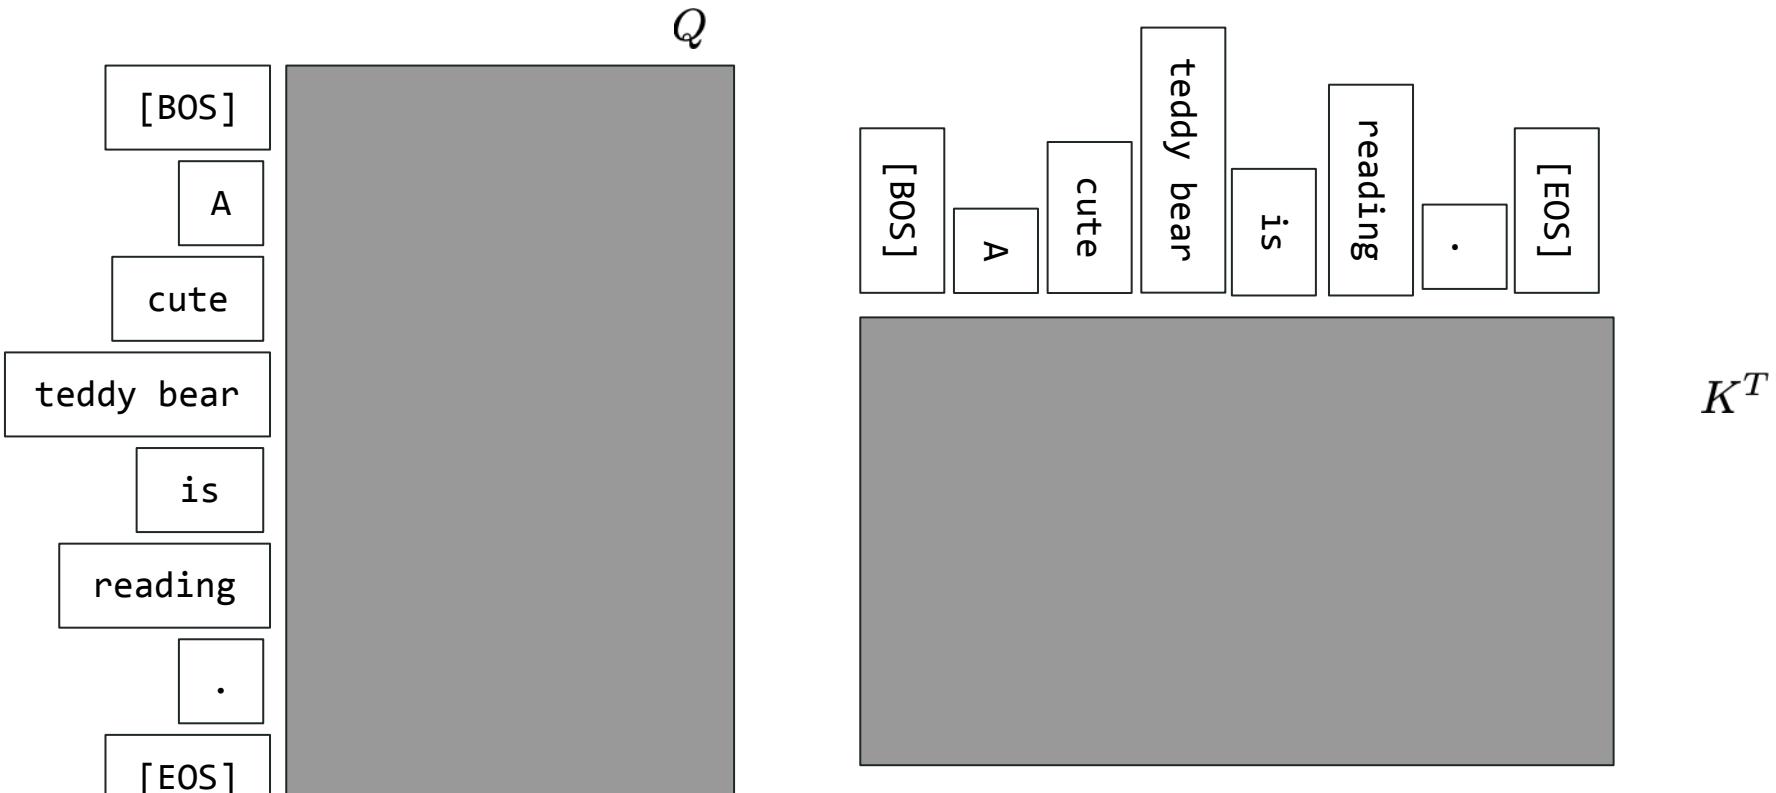
Q



Stitching all the pieces together with an example



Stitching all the pieces together with an example

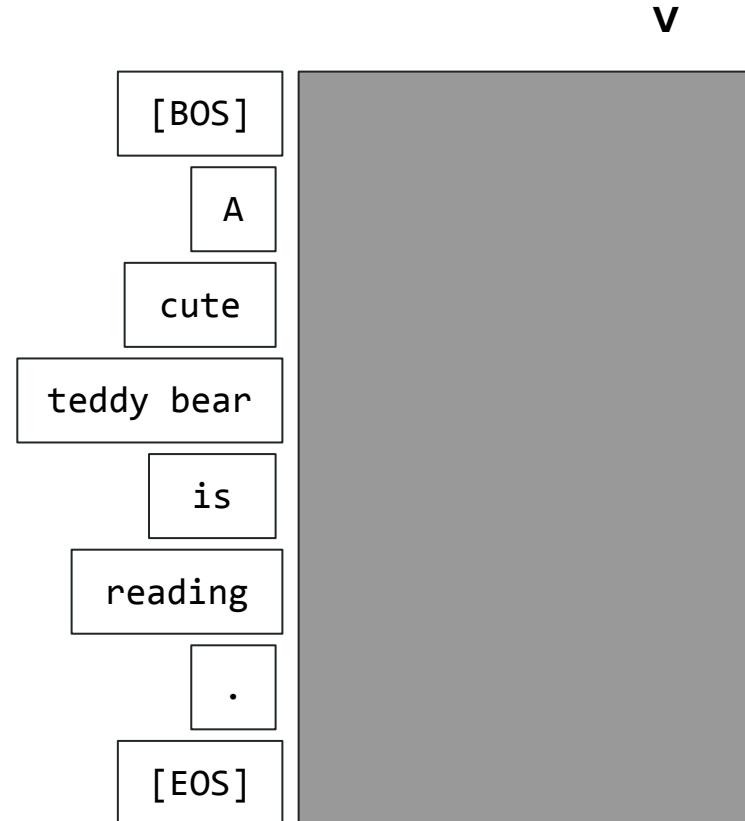
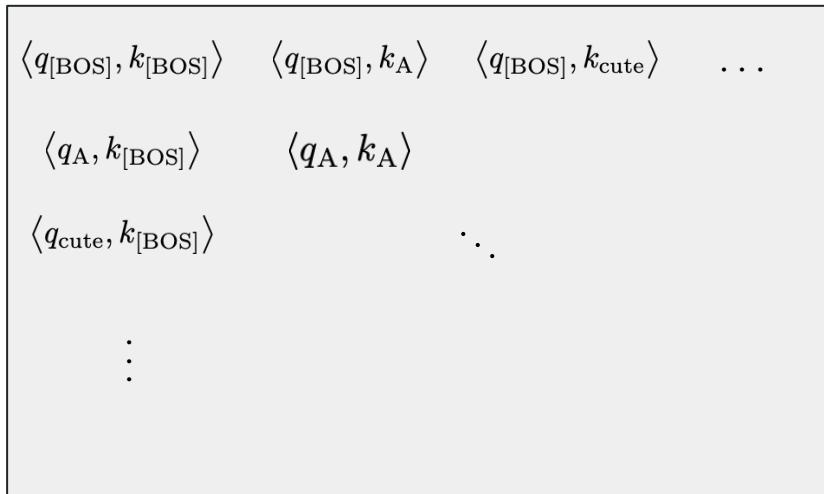


Stitching all the pieces together with an example

$$\langle q_{[\text{BOS}]}, k_{[\text{BOS}]} \rangle \quad \langle q_{[\text{BOS}]}, k_A \rangle \quad \langle q_{[\text{BOS}]}, k_{\text{cute}} \rangle \quad \dots$$
$$QK^T$$
$$\langle q_A, k_{[\text{BOS}]} \rangle \quad \langle q_A, k_A \rangle$$
$$\langle q_{\text{cute}}, k_{[\text{BOS}]} \rangle$$
$$\ddots$$
$$\vdots$$
$$\vdots$$
$$\vdots$$

Stitching all the pieces together with an example

QK^T



Stitching all the pieces together with an example

$QK^T V$

$$\langle q_{[\text{BOS}]}, k_{[\text{BOS}]} \rangle v_{[\text{BOS}]} + \langle q_{[\text{BOS}]}, k_A \rangle v_A + \langle q_{[\text{BOS}]}, k_{\text{cute}} \rangle v_{\text{cute}} + \dots$$

$$\langle q_A, k_{[\text{BOS}]} \rangle v_{[\text{BOS}]} + \langle q_A, k_A \rangle v_A + \langle q_A, k_{\text{cute}} \rangle v_{\text{cute}} + \dots$$

•
•
•

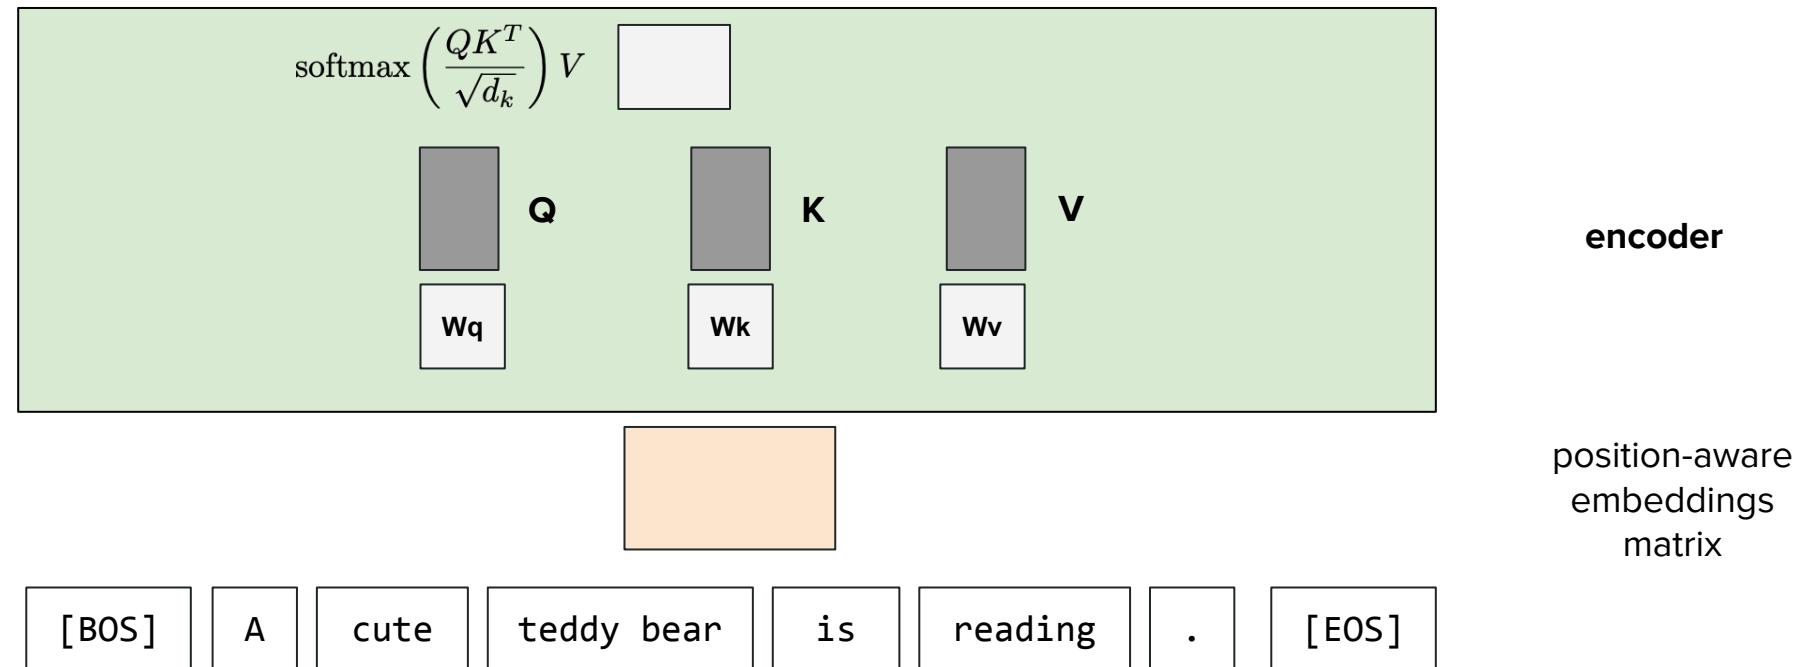
Stitching all the pieces together with an example

$$\text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

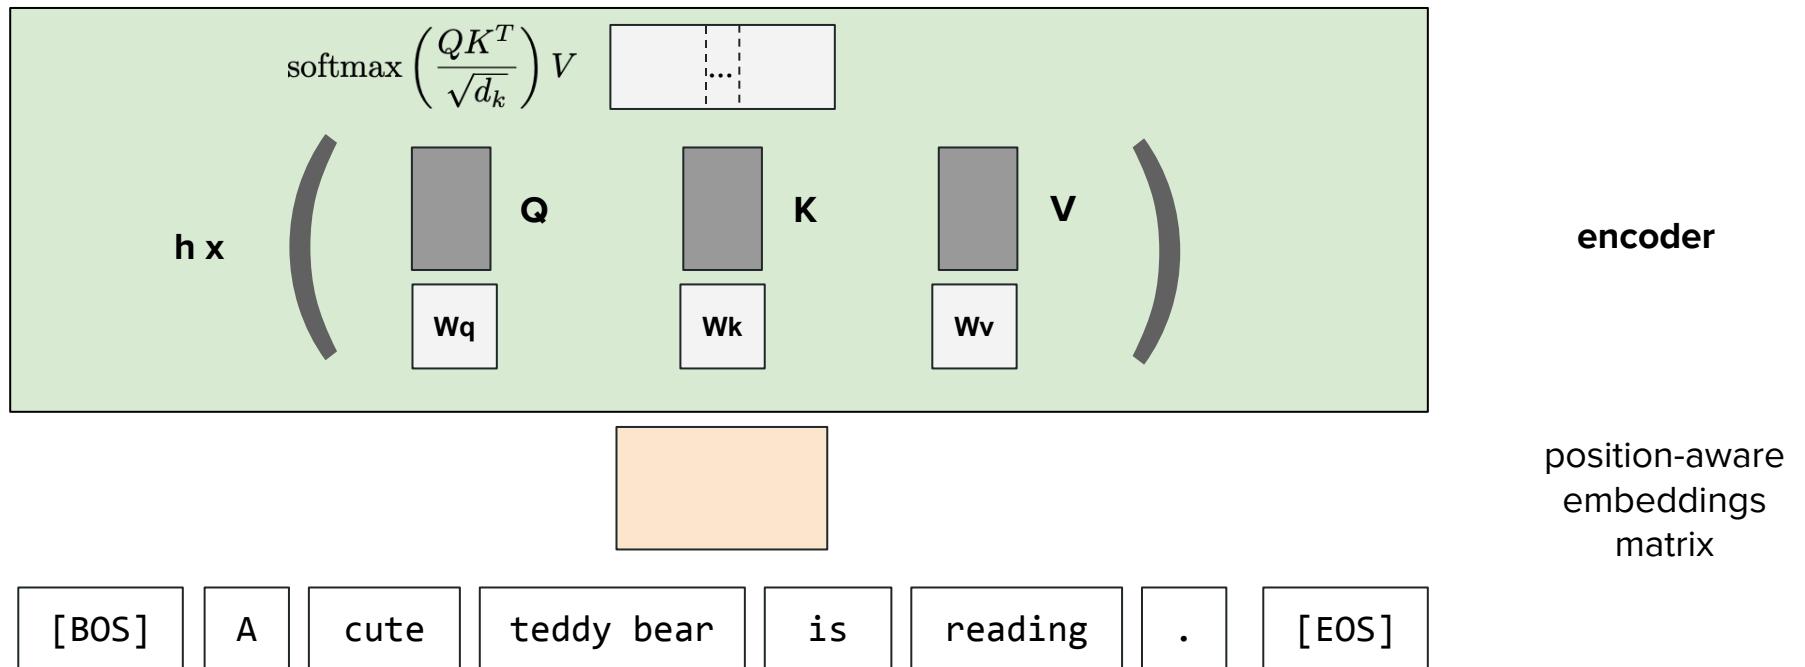
weighted average of values

with weights being a function of $\langle q, k \rangle$

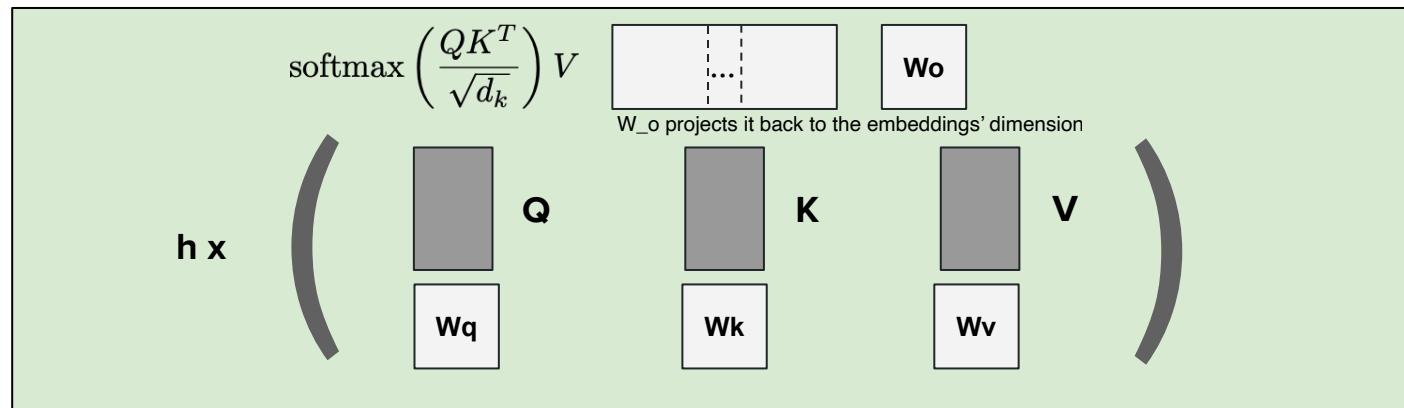
Stitching all the pieces together with an example



Stitching all the pieces together with an example



Stitching all the pieces together with an example



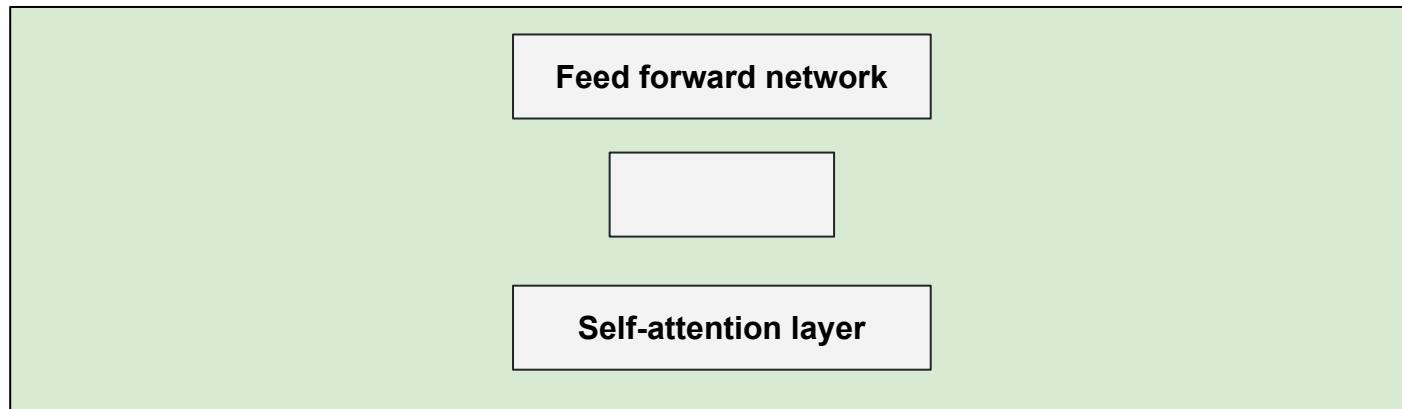
encoder

position-aware
embeddings
matrix



Stitching all the pieces together with an example

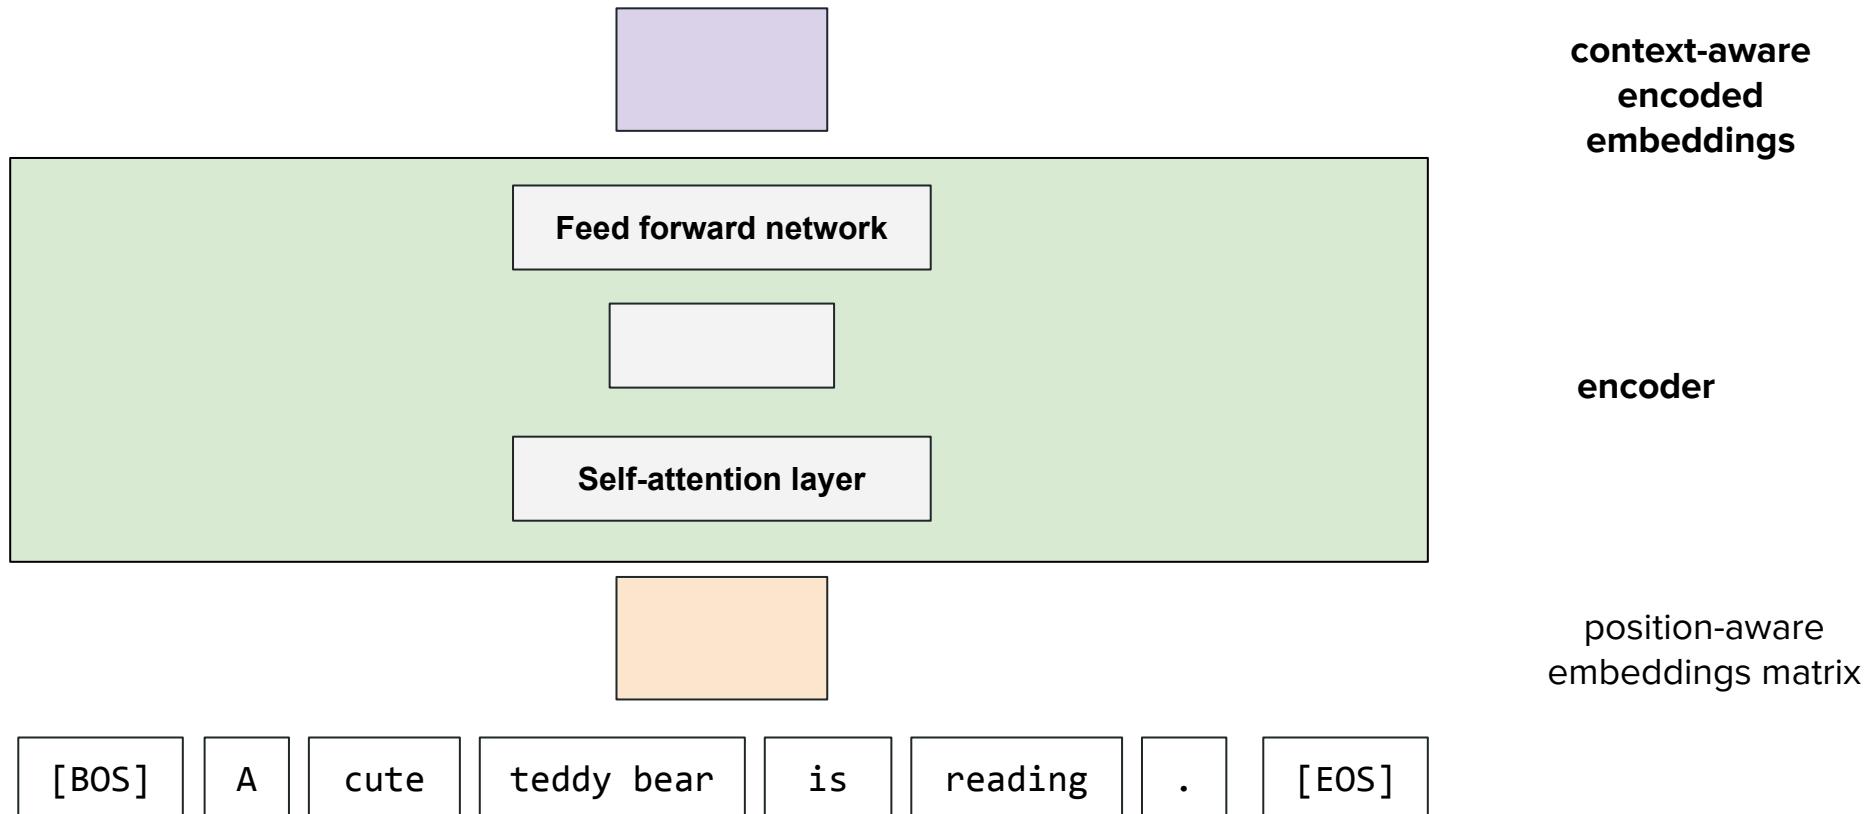
here hidden layer is of bigger dimension than input and output



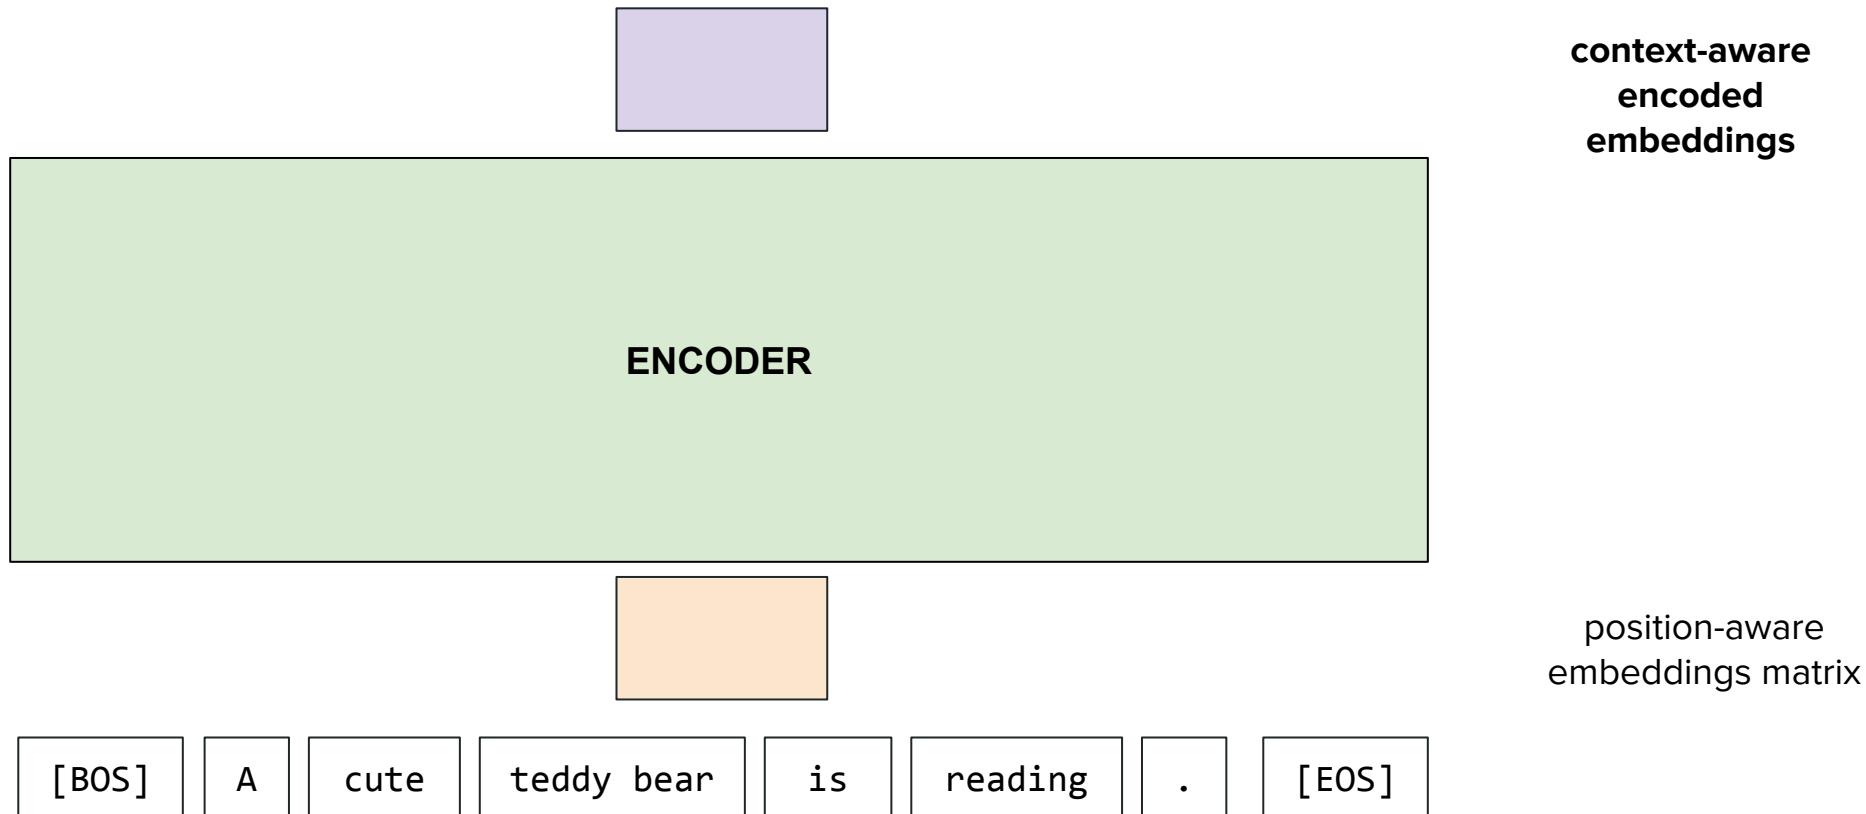
encoder

position-aware
embeddings matrix

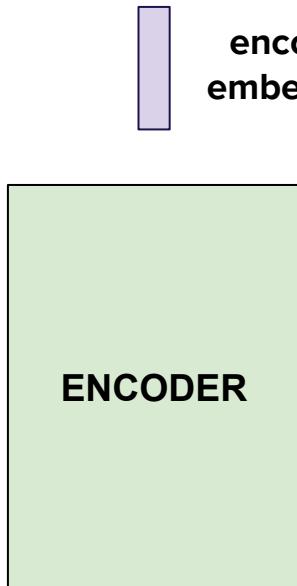
Stitching all the pieces together with an example



Stitching all the pieces together with an example

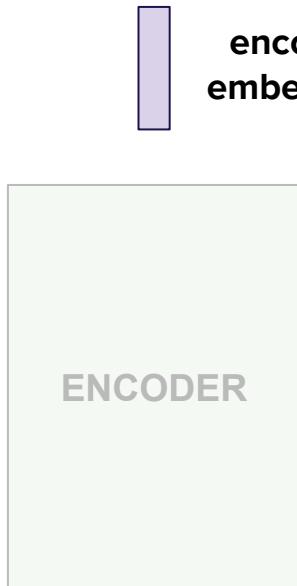


Stitching all the pieces together with an example

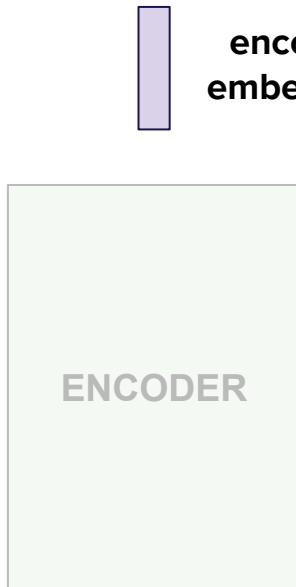


A cute teddy bear
is reading.

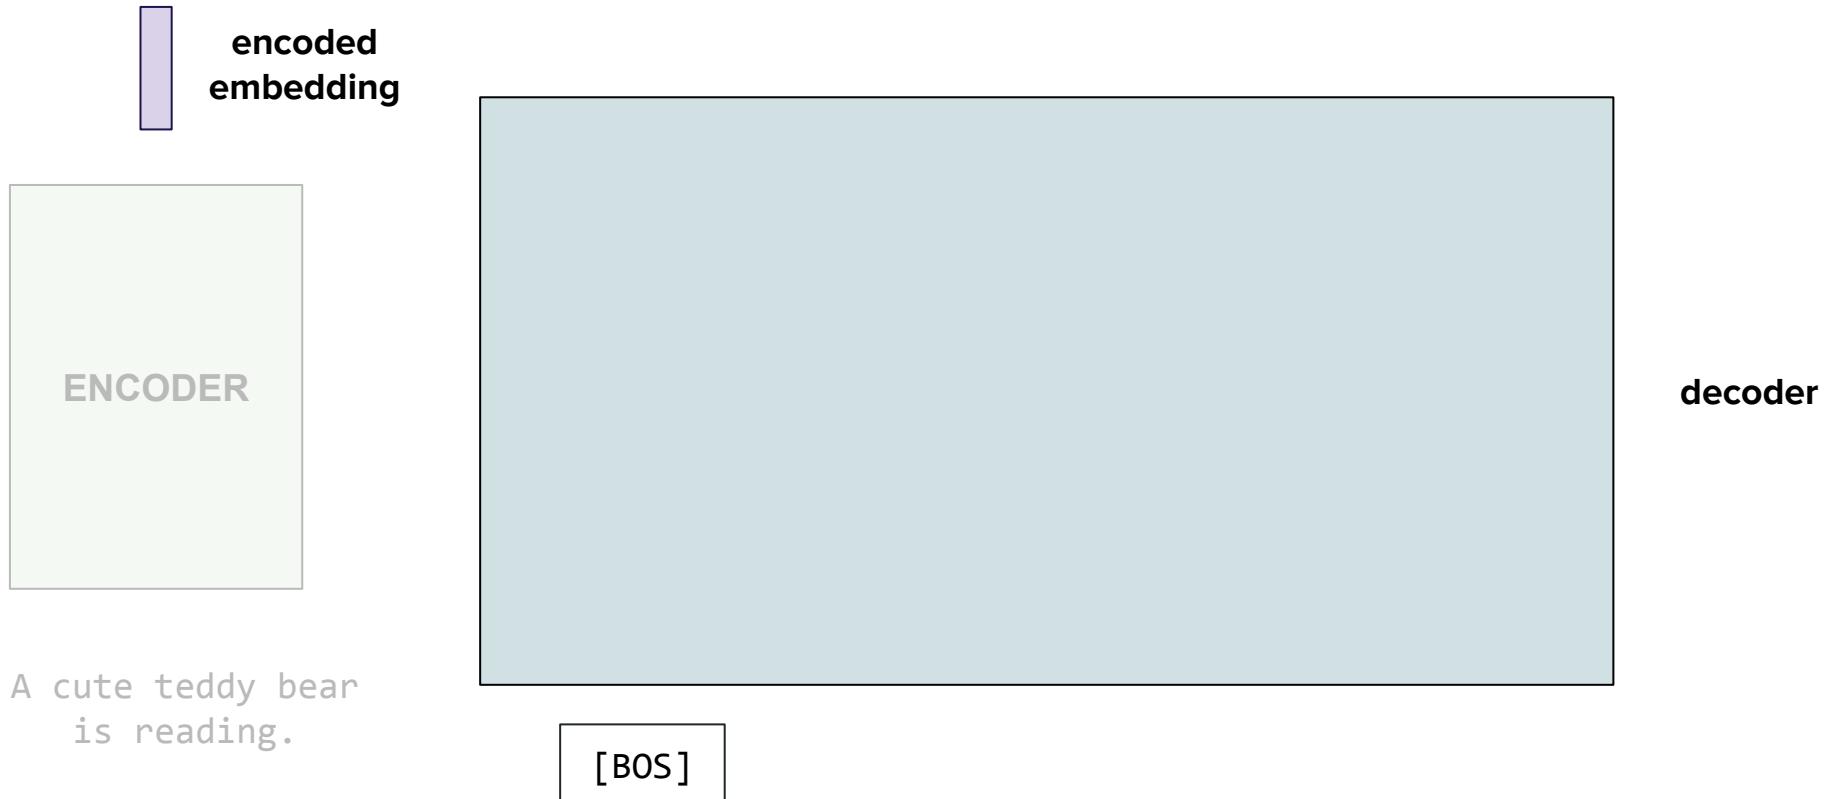
Stitching all the pieces together with an example



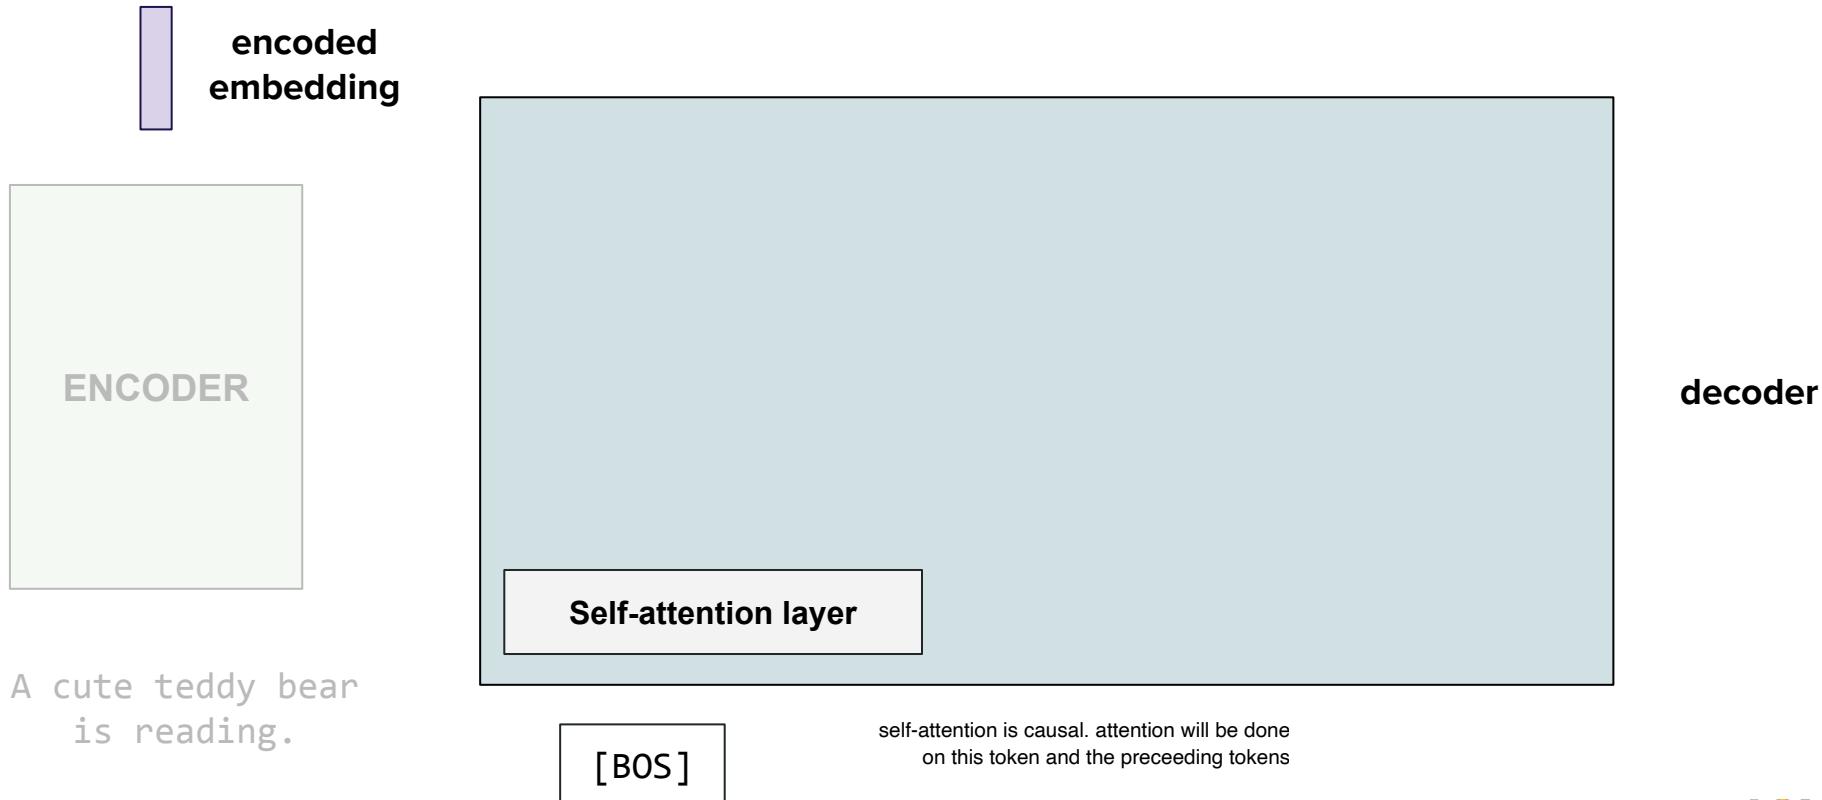
Stitching all the pieces together with an example



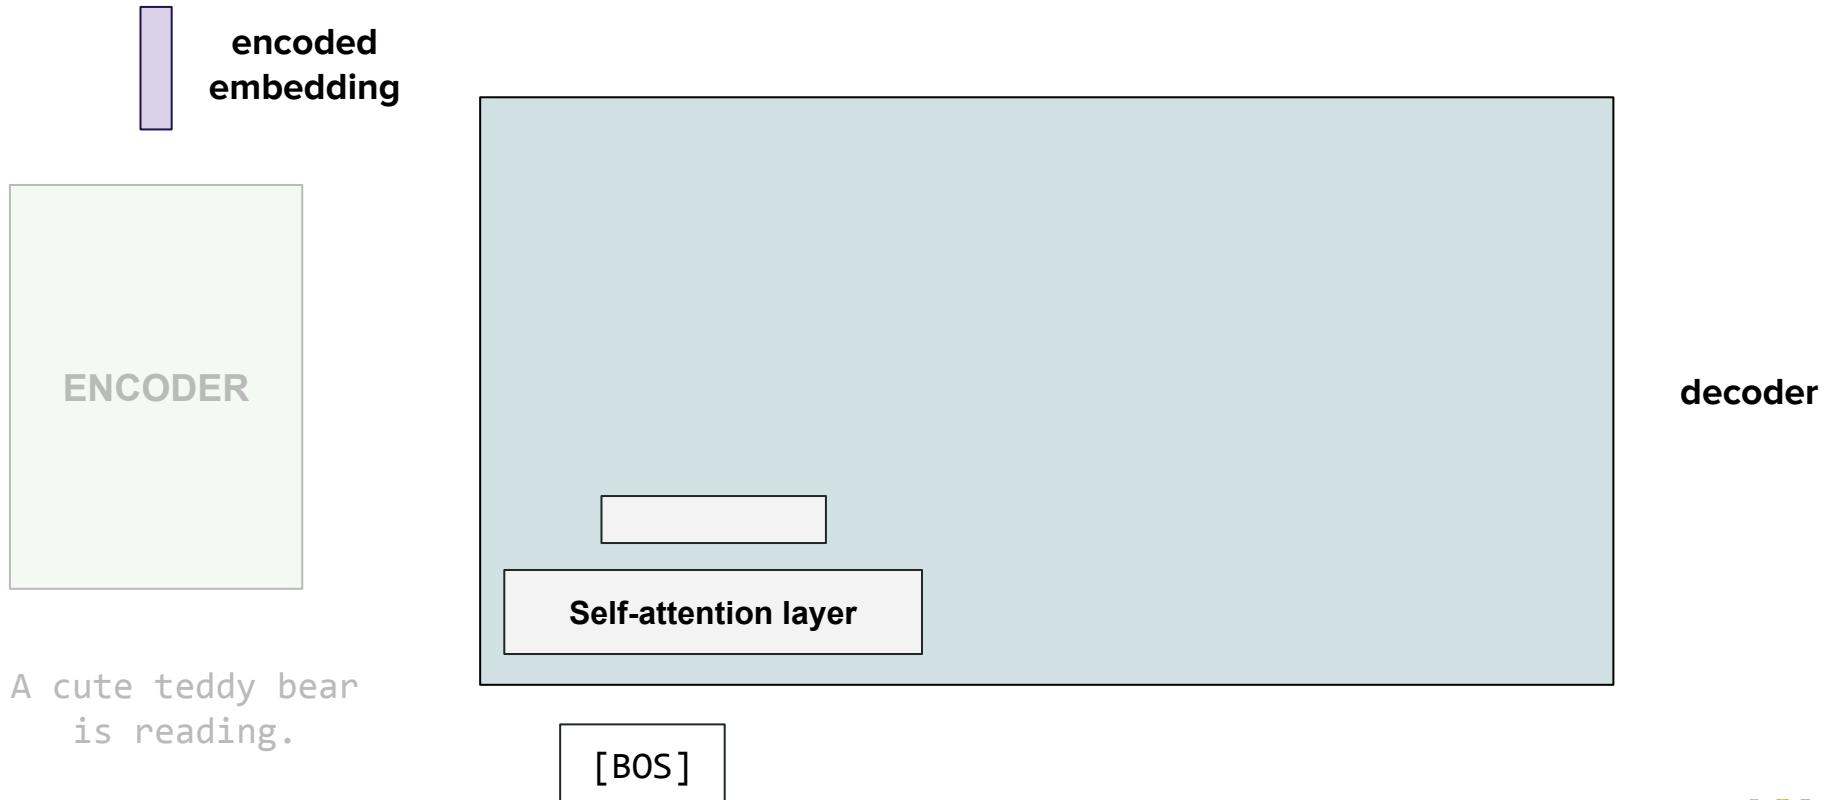
Stitching all the pieces together with an example



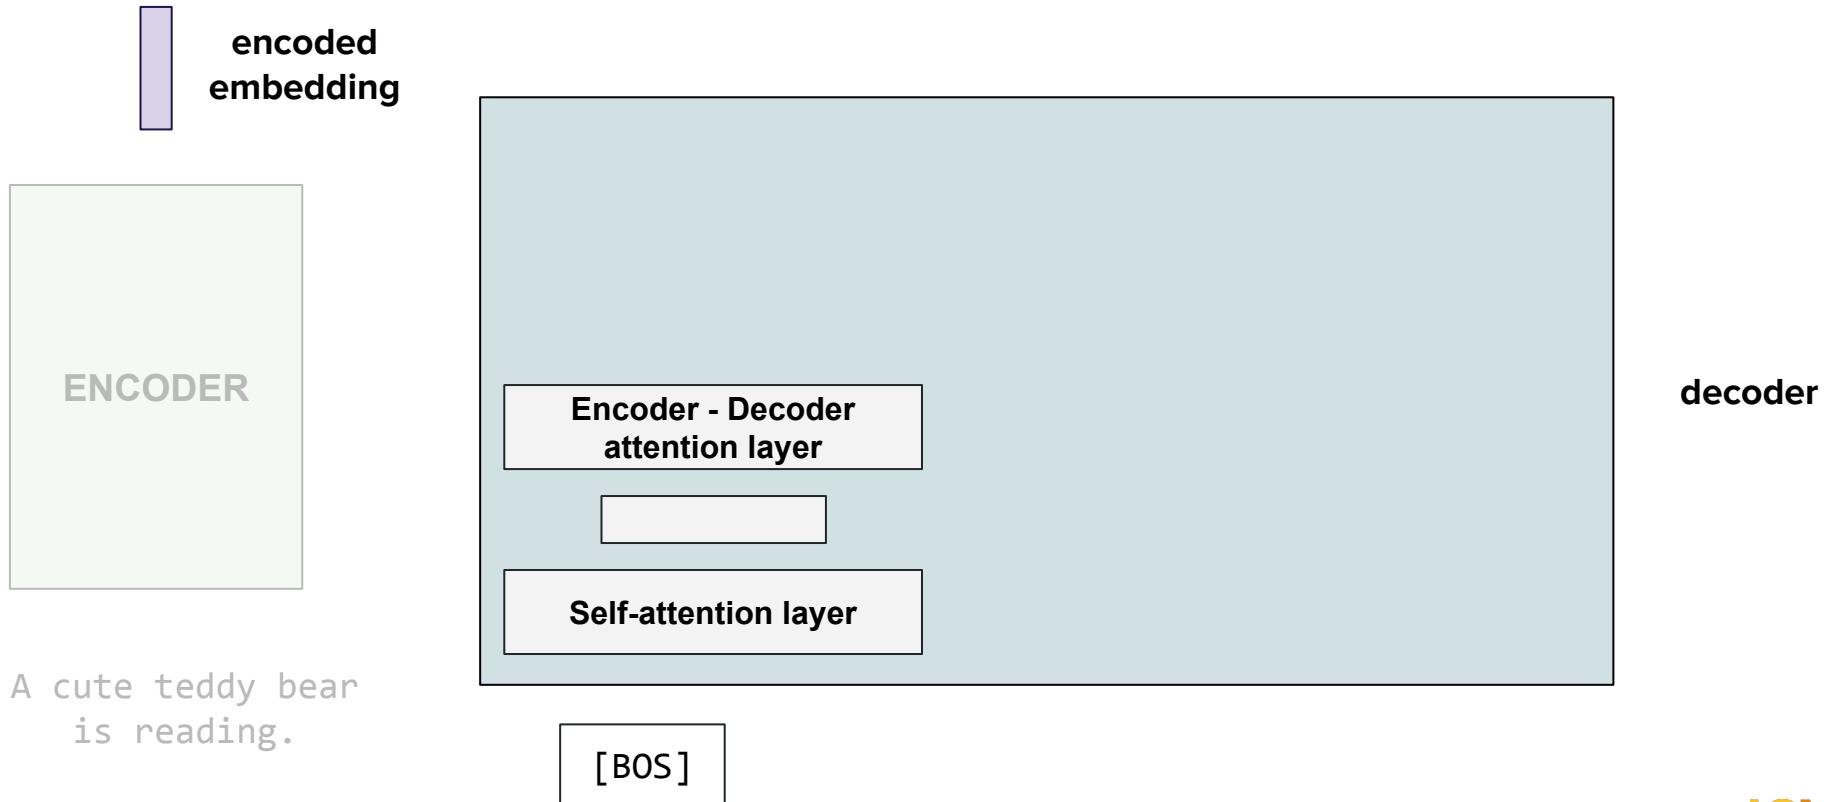
Stitching all the pieces together with an example



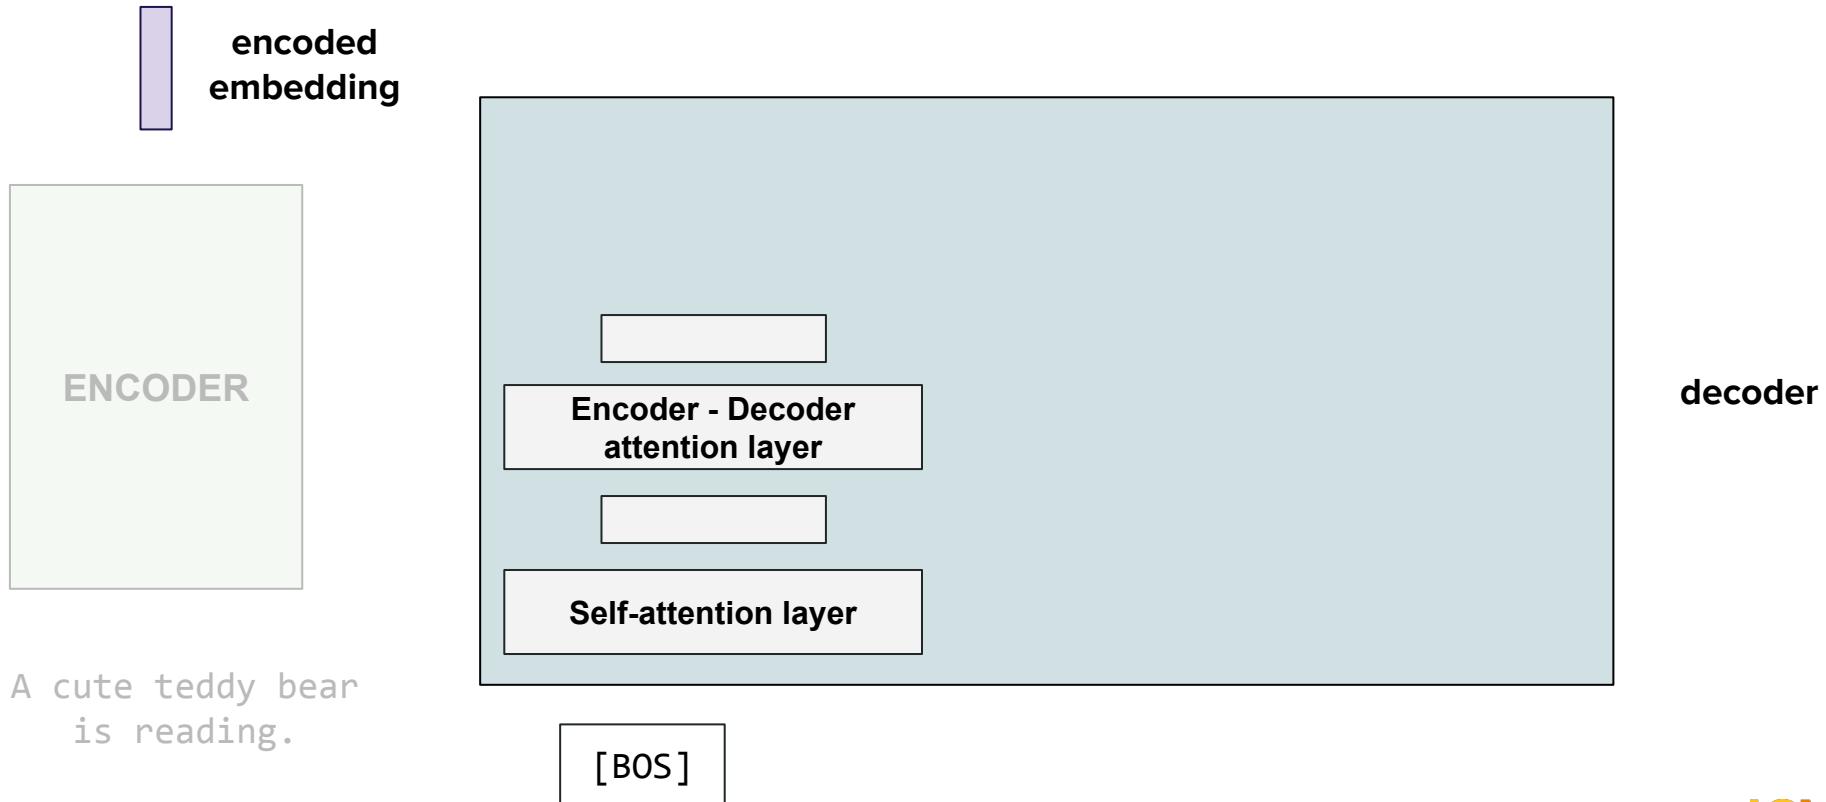
Stitching all the pieces together with an example



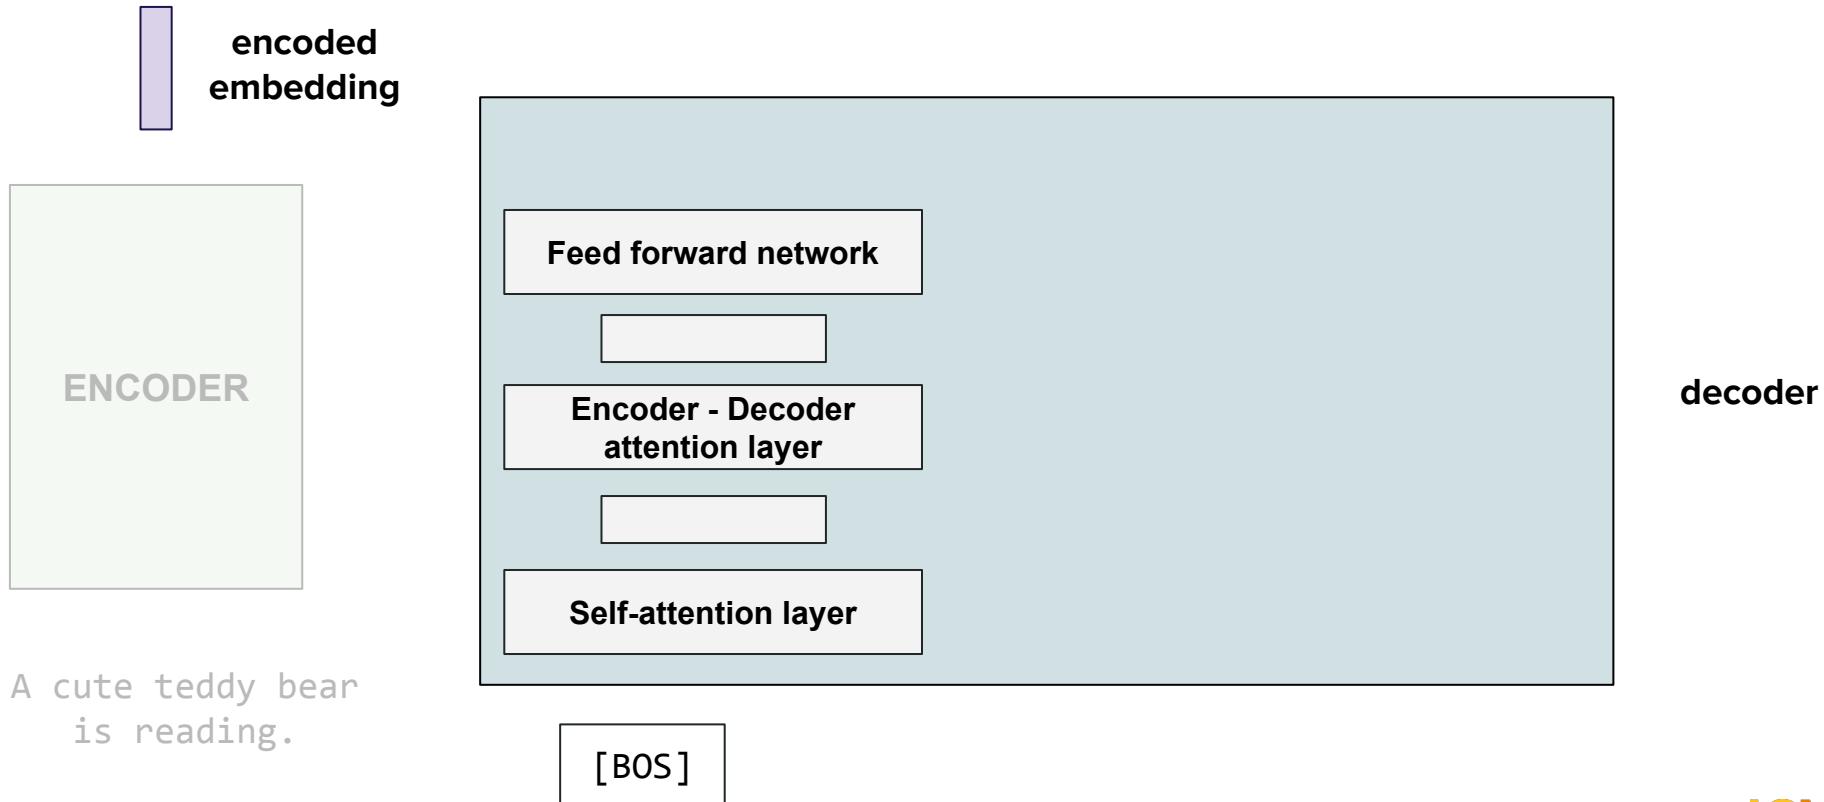
Stitching all the pieces together with an example



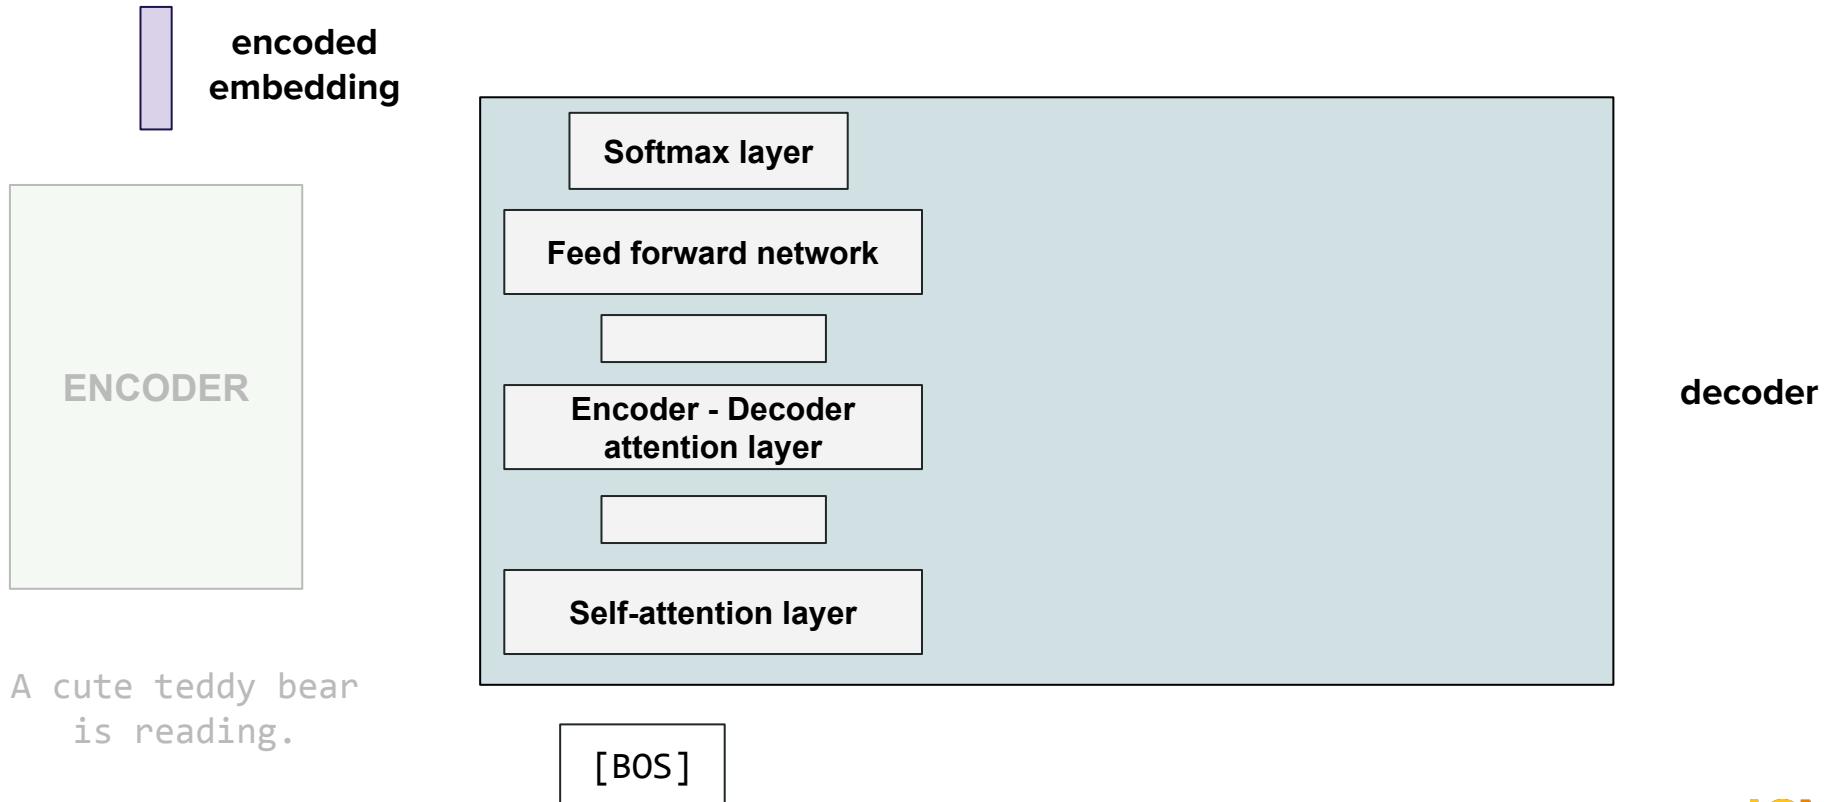
Stitching all the pieces together with an example



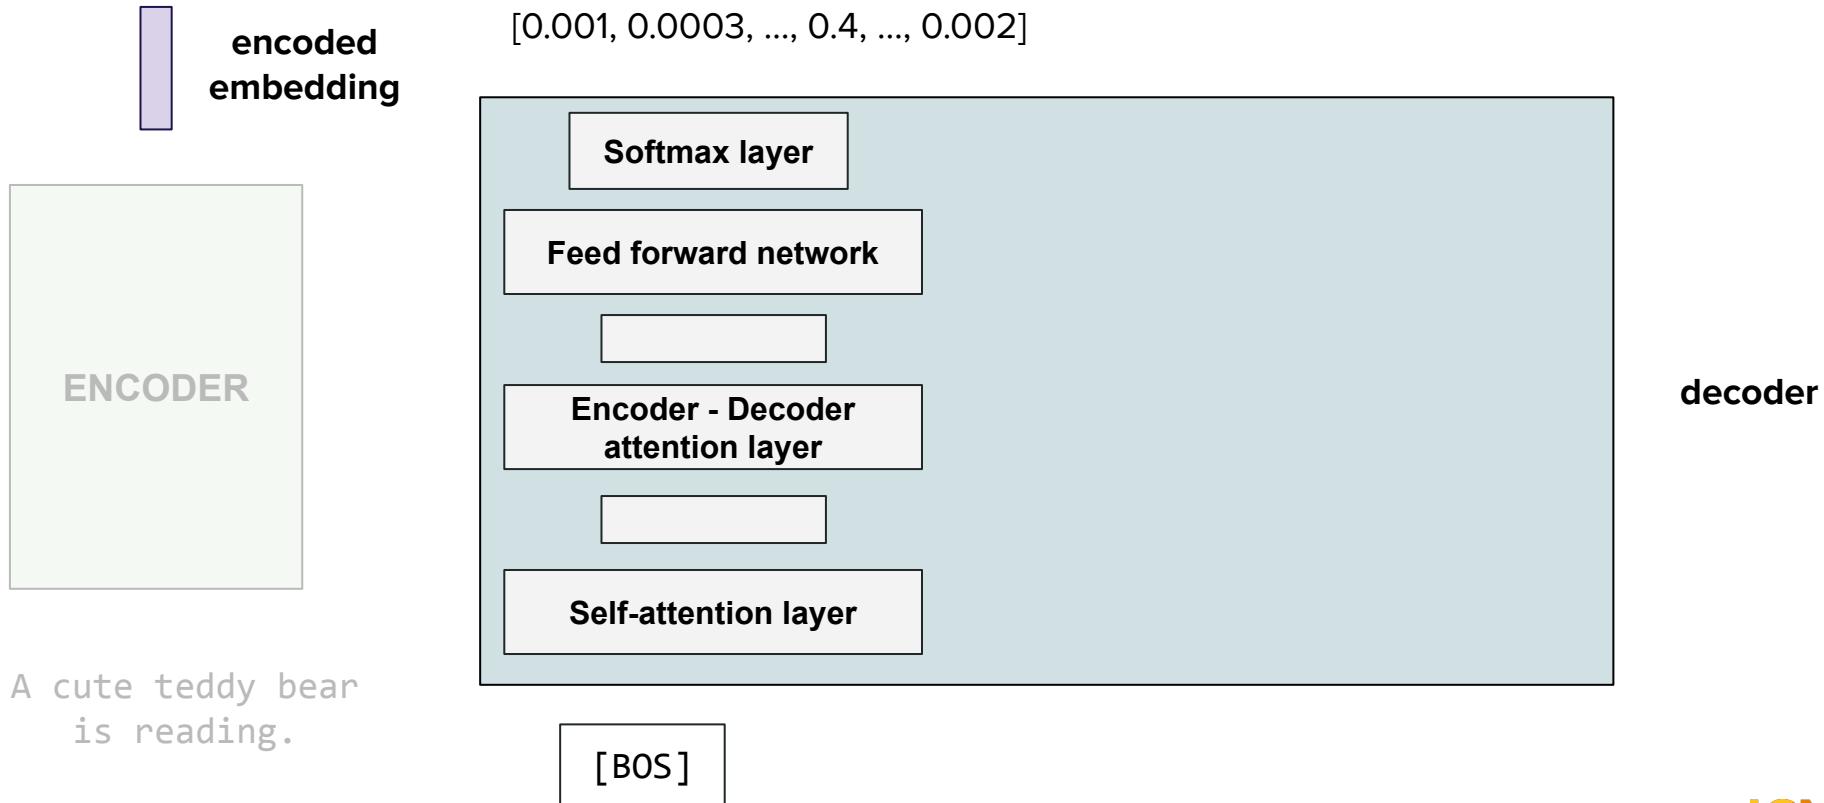
Stitching all the pieces together with an example



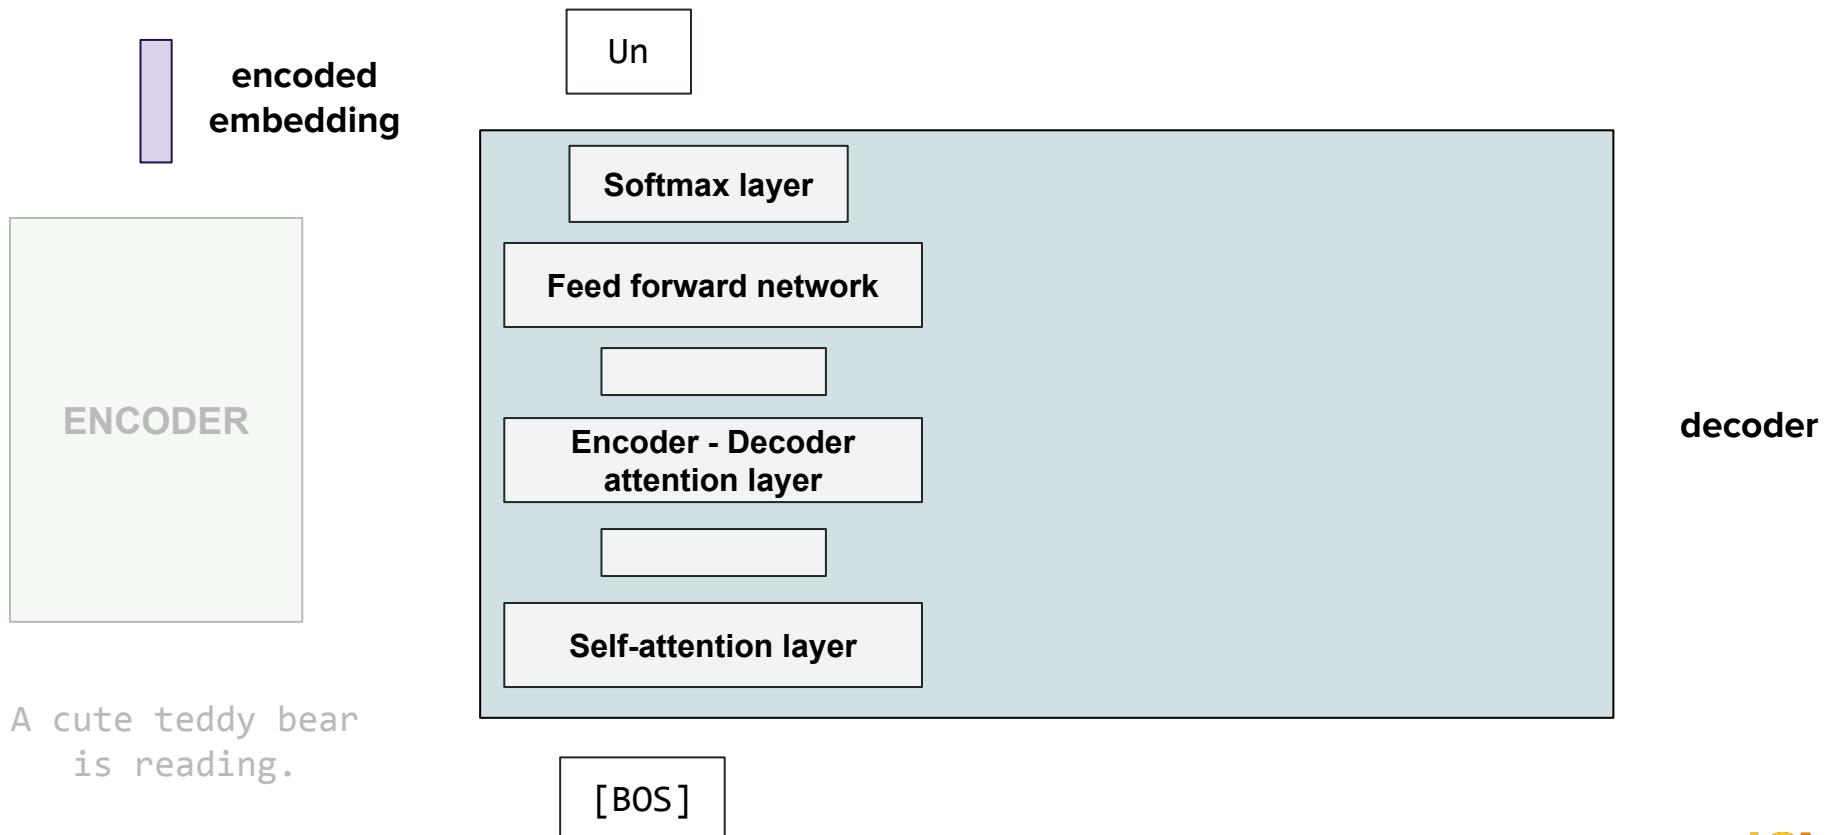
Stitching all the pieces together with an example



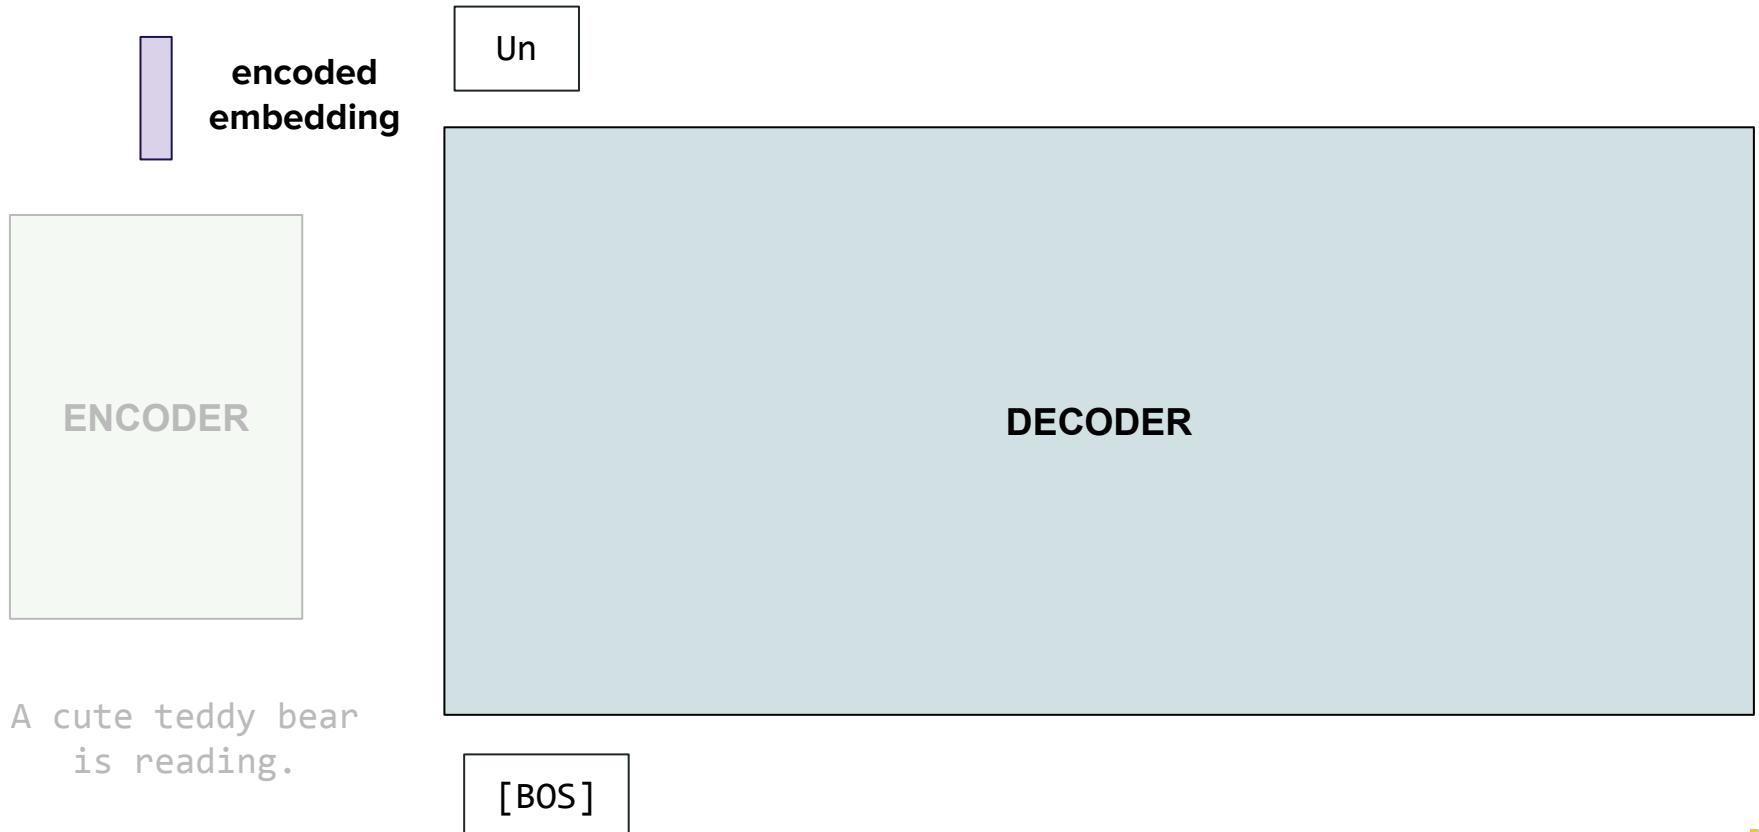
Stitching all the pieces together with an example



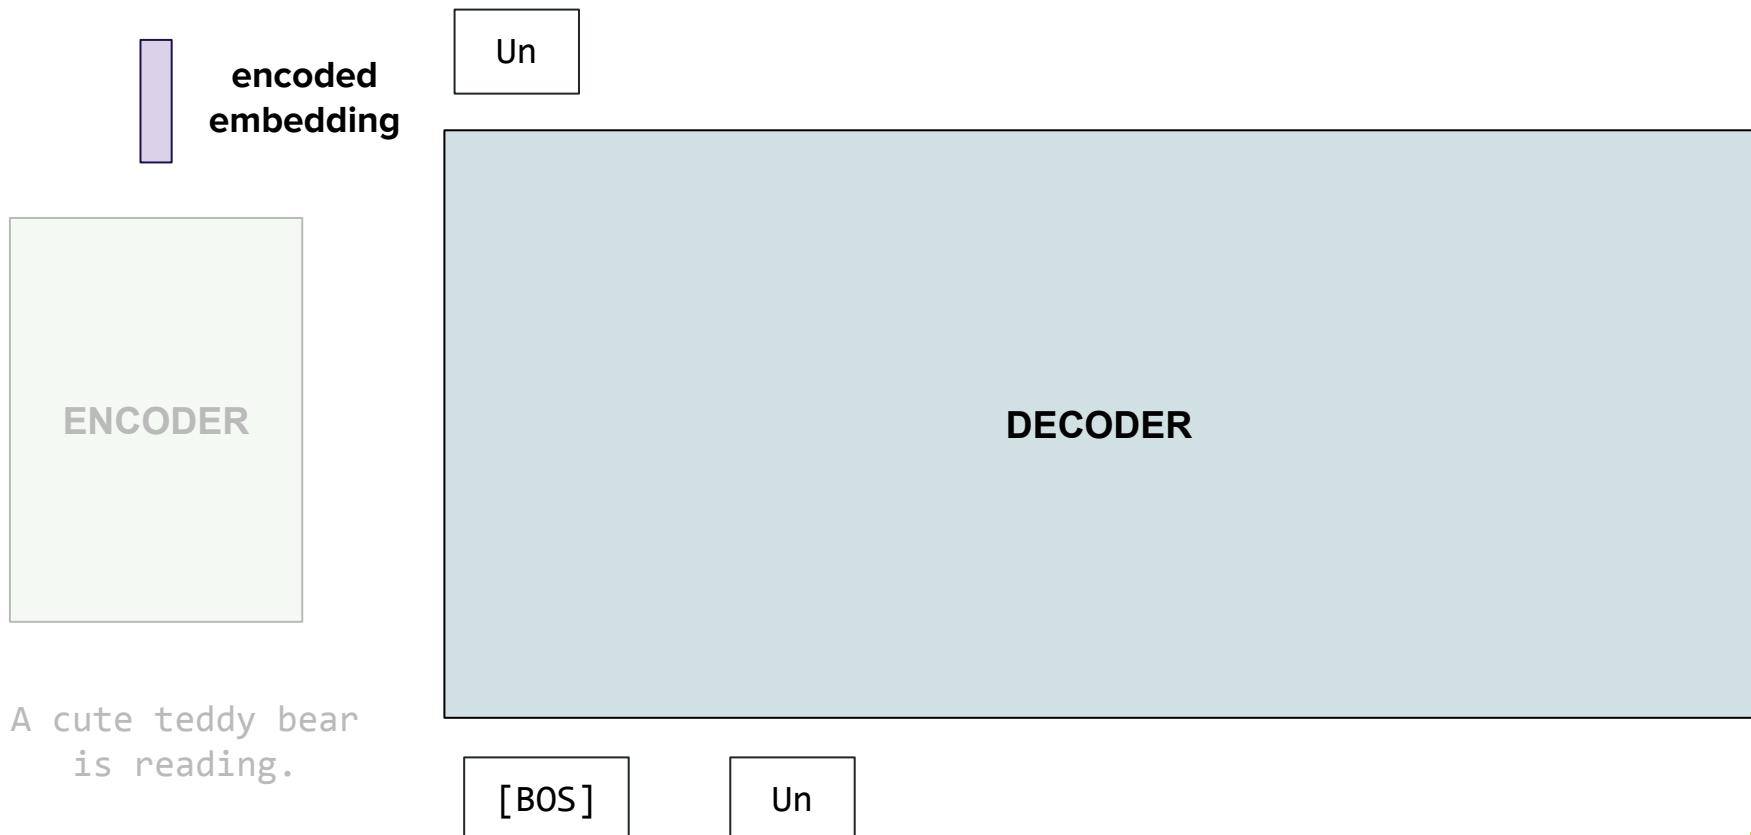
Stitching all the pieces together with an example



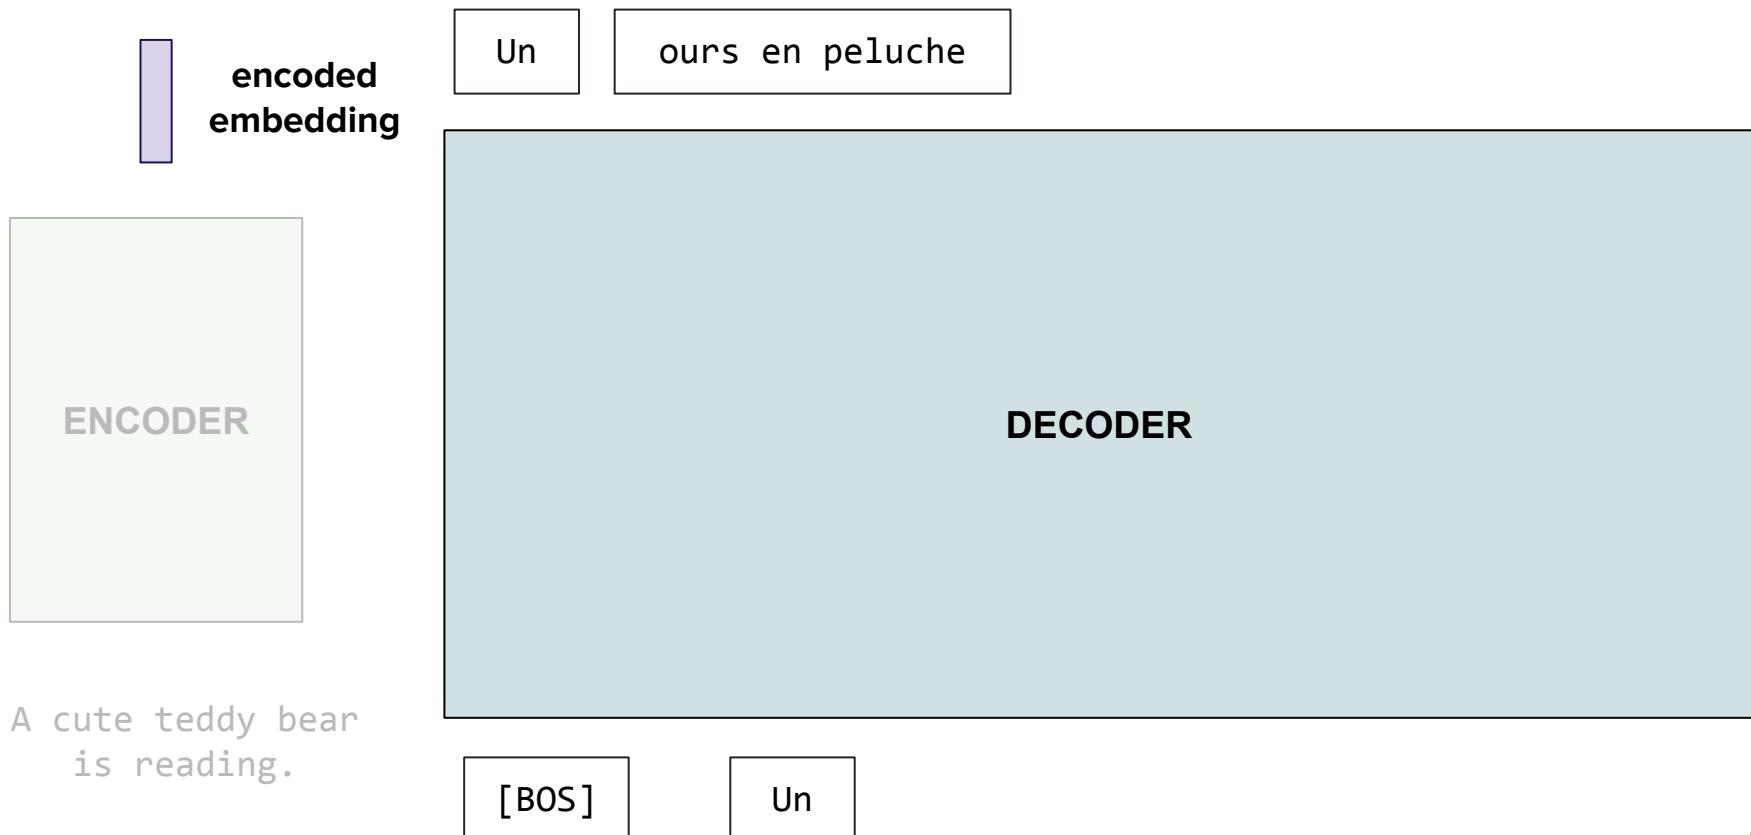
Stitching all the pieces together with an example



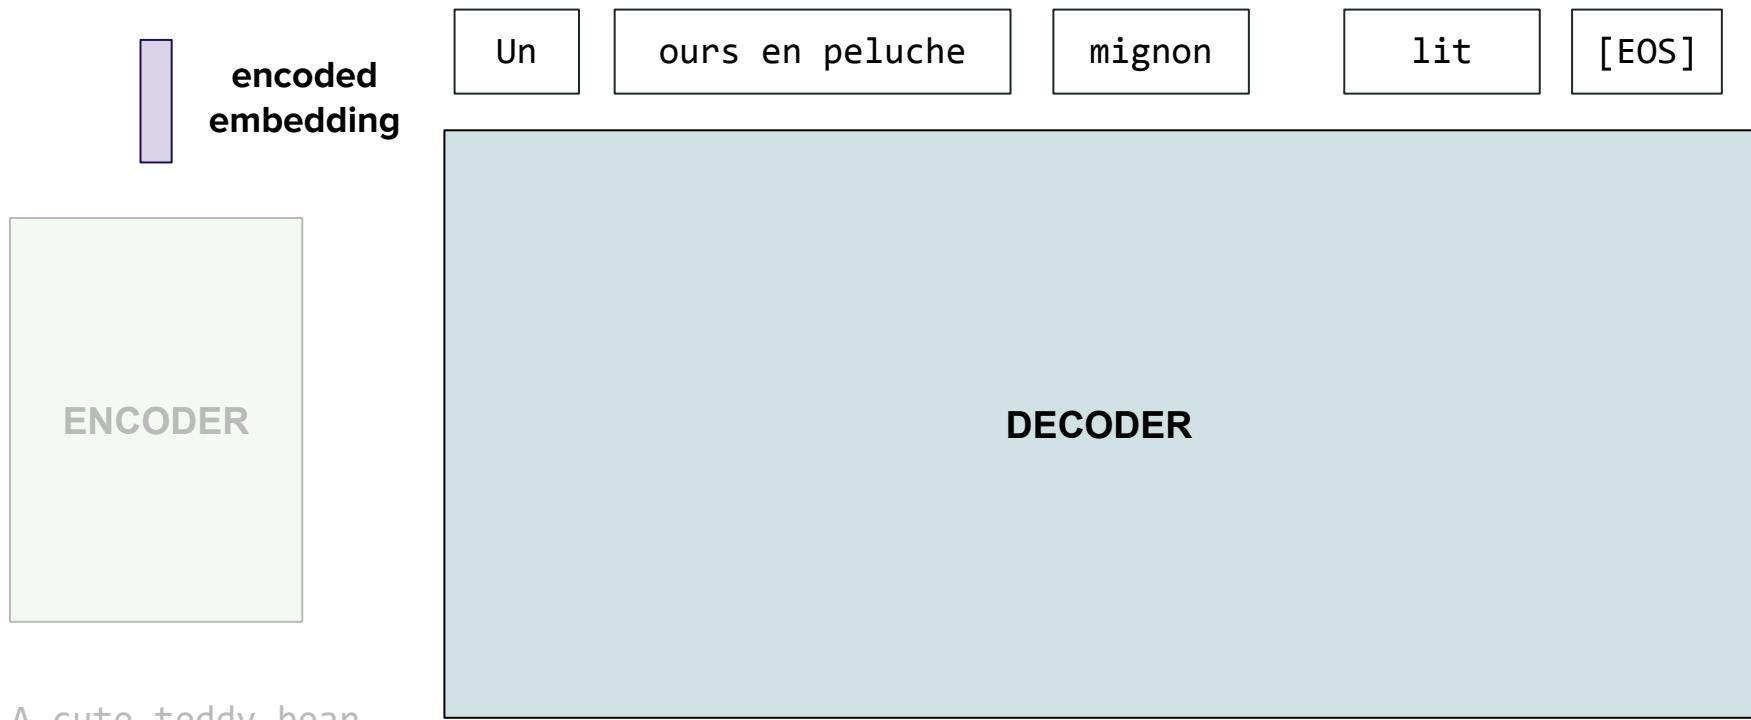
Stitching all the pieces together with an example



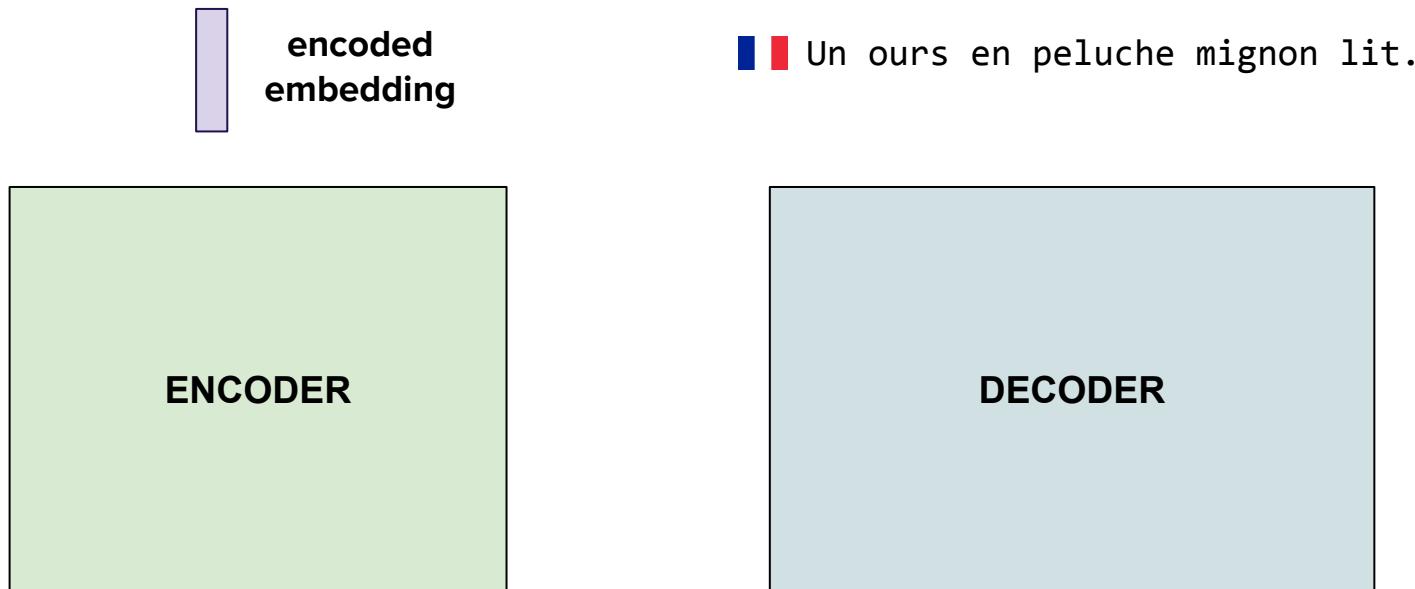
Stitching all the pieces together with an example



Stitching all the pieces together with an example



Stitching all the pieces together with an example



🇺🇸 A cute teddy bear is reading.

Thank you for your attention!
